

Distributed Systems

COMP 212

Lectures 1 and 2

Othon Michail

Course Information

- Lecturer:
 - Othon Michail
 - Office 2.14 Holt Building
- Module Website:
 - **VITAL** 201920-COMP212-DISTRIBUTED SYSTEMS will be used for all module material, assessment submission and feedback
- Structure
 - 30 Lectures + 10 lab practicals
- Assessment
 - **2 programming assignments** (10% each) and **final exam** (80%)

Recommended Reading

- Main Books:
 - Maarten van Steen, Andrew S. Tanenbaum. Distributed Systems, 3rd edition, 2017 (free digital copy: <https://www.distributed-systems.net/index.php/books/ds3/>)
 - Andrew S. Tanenbaum, Maarten van Steen. Distributed Systems: Principles and Paradigms, Pearson, 2nd edition, 2006
 - Nancy Lynch. Distributed Algorithms, Morgan Kaufmann Publishers, 1996
 - Hagit Attiya and Jennifer Welch, Distributed Computing Fundamentals, Simulations, and Advanced Topics, 2nd Edition, Wiley-Interscience, 2004
- Java references
 - Thinking in Java
 - <https://www.mindviewllc.com/quicklinks/>
 - The Java Tutorials (Oracle)
 - <https://docs.oracle.com/javase/tutorial/>

Other Relevant Books

- George Coulouris, Jean Dollimore, Tim Kindberg, Gordon Blair.
Distributed Systems: Concepts and Design, Pearson, 5th Edition, 2011
- Marko Boger. Java in Distributed Systems: Concurrency, Distribution and Persistence, John Wiley & Sons, 1st Edition, 2001

Content of the Course

- Introduction to Distributed Systems

Part I: Distributed Algorithms

- Broadcast (2 lectures)
- Leader Election (3 lectures)
- BFS

Part II: Distributed Systems (and some more algorithms)

- Communication (4 lectures)
- Architectures and Processes
- Naming
- Synchronisation
- Fault tolerance (2 lectures including some algorithms)
- Transactions
- Concurrency and Mutual Exclusion (includes some algorithms)
- Consistency and Replication
- Distribution Protocols
- Security
- Virtualisation and Cloud Computing

Additionally

- Research Lecture

Aims

- Provide an understanding of the technical issues involved in the design of modern distributed systems
- Present some of the major current paradigms

Learning Outcomes

- Appreciation of the **main principles** underlying distributed systems:
 - processes, communication, naming, synchronisation, consistency, fault tolerance, and security
- Familiarity with some of the **main paradigms** in distributed systems
- Knowledge and understanding of the essential facts, concepts, **principles and theories** related to distributed computing
- In depth understanding of the appropriate theory, practices, languages and tools for the specification, design, implementation and evaluation of distributed systems


Introduction

What is a Distributed System?

- A distributed system is:

A collection of independent computers that appears to its users as a single coherent system

What is a Distributed System?



A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.

— *Leslie Lamport* —

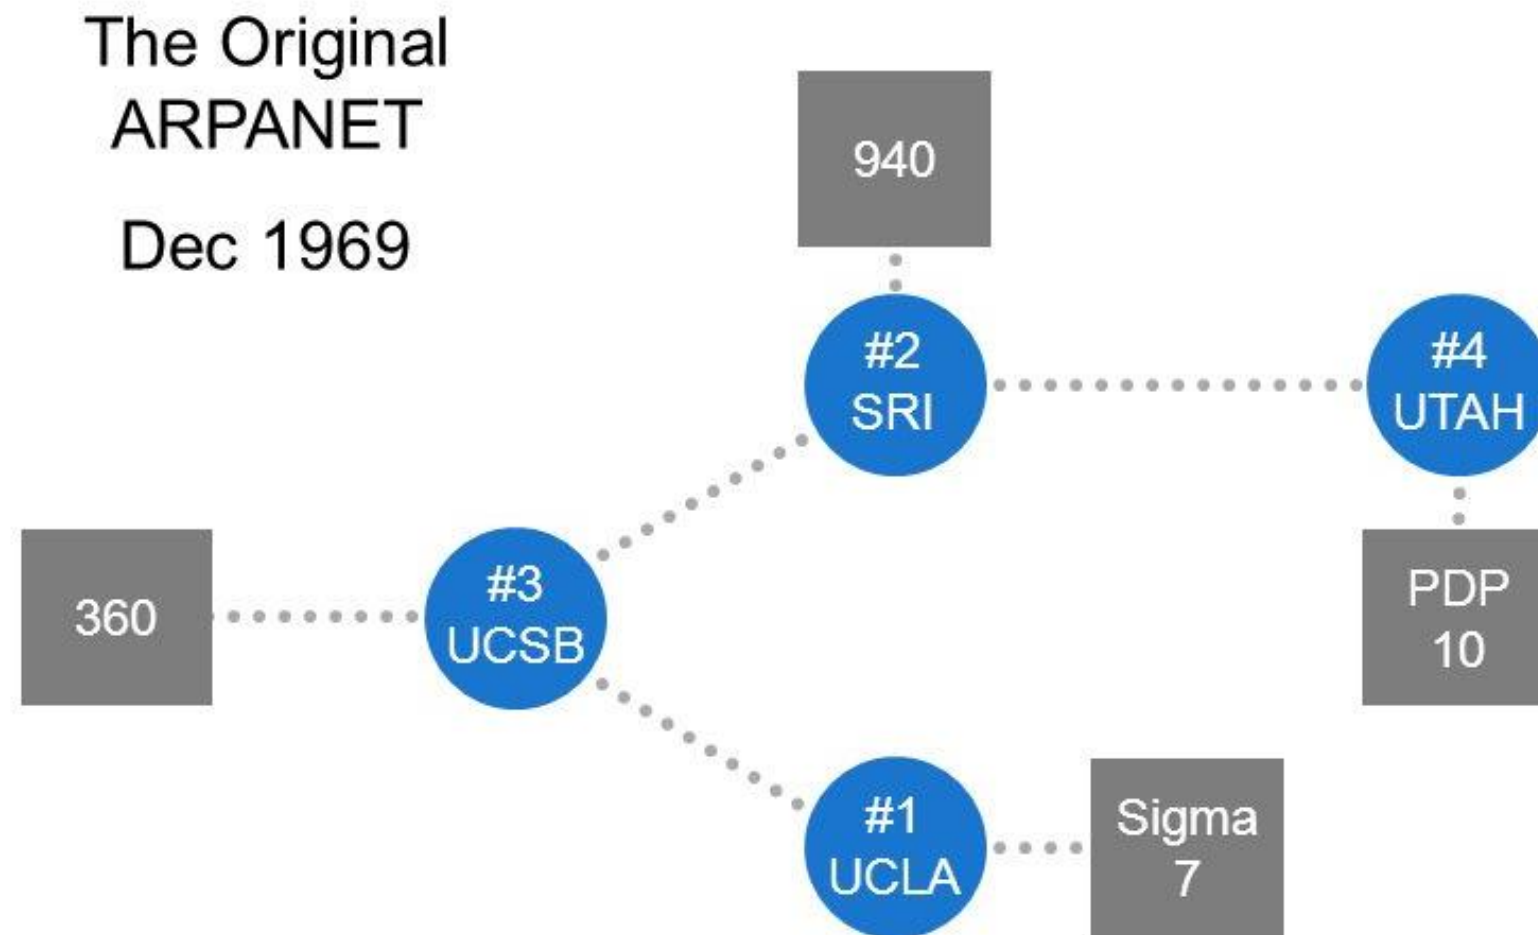
Examples of Distributed Systems

- Can you think of some examples?

Examples of Distributed Systems


- University computer network
- The Internet
- Real-time process control
 - Aircraft/Industrial control systems
- Distributed databases
- Distributed file systems
- Distributed information processing systems
 - Banks (Cash machines)
 - Ticket reservation systems
- Parallel computing (cluster computing, multiprocessor systems)
- Cloud computing

The Beginning of the Internet (ARPANET)



- 4 interconnected computers forming a simple packet switching network
- First to implement TCP/IP


The Internet



- Over 1.5 billion websites
- Over 4 billion users
- Check the rate of growth at:
<http://www.internetlivestats.com/>

Figure 1


Global Internet users and penetration rate (1995-2009)




Sources: Nielsen, ITU; A.T. Kearney analysis

Internet hosts 1981-2012





<https://www.isc.org/solutions/survey/history>



Internet-enabled Devices



Nanorobotics and Self-Assembly



Why Distributed Systems?

- **Pervade society**
 - Billions of computers and mobile devices
 - Continuous advances in communication technology
 - Most crucial modern information systems are distributed
- **Turing Award 2013** was awarded to **Leslie Lamport** for his work on Distributed Systems

Goals of Distributed Systems

- Easily connect users/resources
- Transparency
- Openness
- Reliability
- Performance
- Scalability


Beyond Systems

- Distributed systems also require a rigorous theory to support them
- Theory of Distributed Computing
 - The distributed counterpart of the theory of (centralised) computation and of sequential algorithms
- Why?
 - Theory can provide insights into existing systems
 - Allows to argue about *any* potential system
- For example,
 - Prove what is *in principle feasible* in a given setting or
 - Prove *impossibility results* for a given setting



A Formal Model

- n processors
- **Network:** Directed graph $G=(V, E)$
 - $n = |V|$
- **Examples:** ring, tree, complete network
- Distributed algorithms may assume a specific topology or not

Ring Topology



Tree Topology



Processes

- Executed in processors
- A process u :
 - $states_u$ (set of states)
 - initial states $start_u$ and halting states $halt_u$ subsets of $states_u$
 - $msgs_u: states_u \times out-nbrs_u \rightarrow M \cup \{\text{null}\}$
(message-generation function)
 - $trans_u: states_u \times \{M \cup \{\text{null}\}\}^{|in-nbrs_u|} \rightarrow states_u$
(state-transition function)

Synchronous Networks

- All processes initially in initial state(s)
- All processes repeat synchronously (in lock-step):

Step 1

1. Apply message-generation function
2. Messages produced for out-neighbours
3. Messages transmitted through corresponding channels


Step 2

1. Apply state-transition function (taking into account any received messages)
2. Remove all messages from the channels

Step 1 + Step 2 = 1 Round


An Illustration

- $states_u = \{a, b, c\}$, $M = \{m_1, m_2\}$ (also null)
- $msgs_u(a, 1) = m_1$, $msgs_u(a, 2) = null$, $msgs_u(a, 3) = m_2$
- $trans_u(a, (m_1, m_2, null, m_2)) = b$




An Illustration

- $states_u = \{a, b, c\}$, $M = \{m_1, m_2\}$ (also null)
- $msgs_u(a, 1) = m_1$, $msgs_u(a, 2) = null$, $msgs_u(a, 3) = m_2$



An Illustration

- $states_u = \{a, b, c\}$, $M = \{m_1, m_2\}$ (also null)
- $trans_u(a, (m_1, m_2, \text{null}, m_2)) = b$



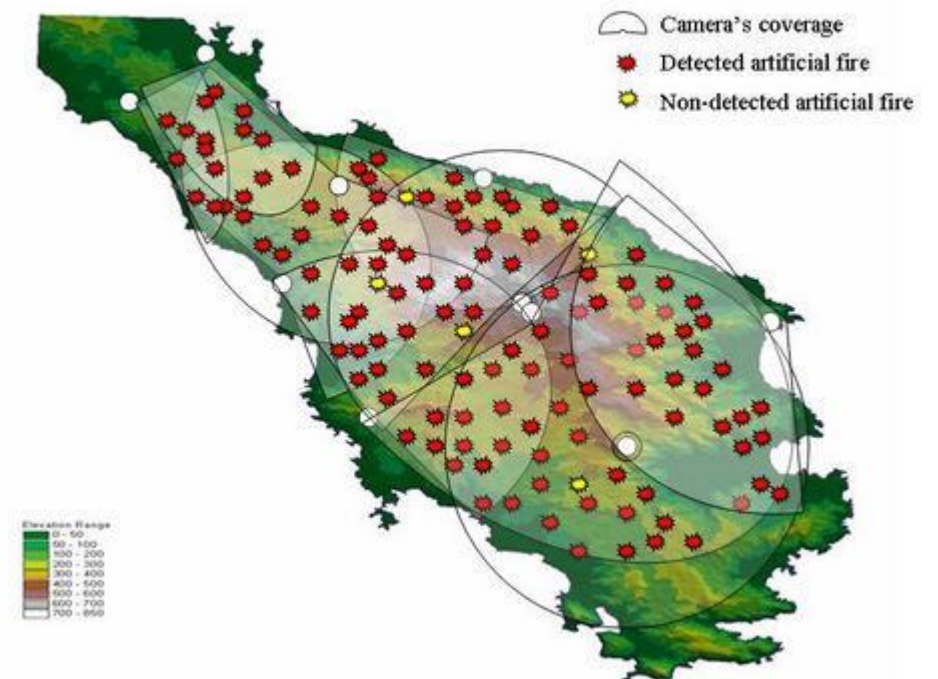
A Simple Example: Flooding

- **Initially:** A **leader-node** has a **token** (piece of info)
- **Goal:** All nodes must learn the token
- **Distributed Algorithm:** Ideas???




Example application:

- Wireless Sensor Network
- A node detects fire
- All nodes must be informed about the location to act accordingly



A Simple Example: Flooding

- Initially: A leader-node has a token (piece of info)
 - Goal: All nodes must learn the token
 - Distributed Algorithm: Ideas???
- 
- Formally, we must define: M , $states_u$, $msgs_u$ and $trans_u$
 - Algorithmic Idea (informal):
 - All nodes awake initially
 - If in the beginning of this round you are awake and you have the token,
 - forward it to all your neighbours and
 - sleep
 - If you are asleep, do nothing