

Distributed Systems

COMP 212


Lecture 4

Othon Michail

Flooding/Broadcast

Broadcast given Spanning Tree

- We start from the case in which a **spanning tree** of the network is given



- Network $G = (V, E)$
- $E' \subseteq E$ specifies a spanning tree $T = (V, E')$
- **Root:** u_0 (**leader**)
- Processors know T in a **distributed way**
- Each u_i knows:
 - a *parent_i*
 - a set *children_i*

Broadcast given Spanning Tree

Problem:

- u_0 has some **information** it wishes to **send to all processors**
 - e.g., a **message $\langle M \rangle$**
 - additionally all nodes must have **terminated** in the end

Solution: Pseudocode

Algorithm **Spanning tree broadcast**

State of processor u_i :

- $parent_i$: holds a processor index or nil; u_i 's parent
- $children_i$: holds a set of processor indices (possibly empty); u_i 's children
- Boolean $terminated_i$: indicates whether u_i has terminated (1) or not (0)

Solution: Pseudocode

Algorithm **Spanning tree broadcast**

Initially u_0 knows $\langle M \rangle$


Code for **leader** (u_0):

- send $\langle M \rangle$ to all children
- terminate


Code for **non-leader**:

- upon receiving $\langle M \rangle$ from parent:
 - send $\langle M \rangle$ to all children
 - terminate

Example Execution




Example Execution




round = 1

Example Execution




round = 1

Example Execution




round = 2

Example Execution




round = 2

Example Execution




round = 3

Example Execution



round = 3

Example Execution



round = 4

Correctness and Performance

When we devise an algorithm we typically should

1. Convince that it is **correct**
2. Analyse its **performance**

- **Correctness:**
 - Usually a proof that the algorithm does as expected
- **Performance:**
 - **Time Complexity** (e.g., #rounds required)
 - **Space Complexity** (e.g., memory used by processors)
 - **Communication Complexity** (e.g., total #messages transmitted, size of messages); also called **message complexity**

Spanning Tree Broadcast Correctness

Fairly easy for this algorithm:

- Show that every node will receive $\langle M \rangle$ and terminate
 - Whenever a u_i receives $\langle M \rangle$ it terminates by the end of that round
 - Suffices to show that every u_i will receive $\langle M \rangle$

Proof. Take any $u_i \neq u_0$. As the tree T is spanning, there is a single tree-path from u_0 to u_i . By the way the algorithm works, $\langle M \rangle$ will be forwarded hop-by-hop on the path until it reaches u_i . And this holds for all u_i in the network. □

S. T. Broadcast Time Complexity

- Equal to the #rounds until all nodes have received $\langle M \rangle$

Lemma. For every u_i whose distance from u_o in the spanning tree is r , it holds that u_i receives $\langle M \rangle$ in round r .

Proof. By induction on r .

- For $r = 1$: Holds because in round 1, u_o transmits $\langle M \rangle$ to all its children who receive it in round 1
- Assume that it holds for any $r - 1 \geq 1$
 - Means that all processors at distance $r - 1$ receive $\langle M \rangle$ in round $r - 1$
- Then it must hold also for r
 - The parent u_i of any u_j at distance r is at distance $r - 1$
 - By previous assumption, the parent received $\langle M \rangle$ in round $r - 1$, therefore transmits it to all its children including u_j in round r , and u_j receives it in round r

□

S. T. Broadcast Time Complexity

- This means that for a tree T of **depth d** the algorithm requires **d rounds**
- But in general we want our algorithm to run on all possible networks and all their possible spanning trees
 - And not all have the same depth...
- What is the **worst-case** time complexity of our algorithm?

S. T. Broadcast Time Complexity

- This means that for a tree T of **depth d** the algorithm requires **d rounds**
- But in general we want our algorithm to run on all possible networks and all their possible spanning trees
 - And not all have the same depth...
- What is the **worst-case** time complexity of our algorithm?
- **It is the worst-case depth of a spanning tree on n nodes**

S. T. Broadcast Time Complexity

- This means that for a tree T of **depth d** the algorithm requires **d rounds**
- But in general we want our algorithm to run on all possible networks and all their possible spanning trees
 - And not all have the same depth...
- What is the **worst-case** time complexity of our algorithm?
- It is the worst-case depth of a spanning tree on n nodes **$= n - 1$**

S. T. Broadcast Communication Complexity

1. Size of largest message transmitted?

S. T. Broadcast Communication Complexity

1. Size of largest message transmitted?
 - Size of $\langle M \rangle$, typically in bits
- For example, if $\langle M \rangle$ is a single **processor identifier** (or *id*) this would typically be $O(\log n)$ bits

S. T. Broadcast Communication Complexity

2. Total **#messages** transmitted?

- A single observation suffices
- **Any ideas?**

S. T. Broadcast Communication Complexity

2. Total #messages transmitted?

- A single observation suffices

Observation. *For every edge of the spanning tree, from a parent u_i to a child u_j , exactly one message will be ever transmitted through it.*

- The single transmission of $\langle M \rangle$ from u_i to u_j
- As u_i then terminates it will not happen again
- What is the total #messages then?

S. T. Broadcast Communication Complexity

2. Total #messages transmitted?

- A single observation suffices

Observation. *For every edge of the spanning tree, from a parent u_i to a child u_j , exactly one message will be ever transmitted through it.*

- The single transmission of $\langle M \rangle$ from u_i to u_j
- As u_i then terminates it will not happen again
- What is the total #messages then?
 - #edges of a spanning tree on n nodes
 - Always $n - 1$

Spanning Tree Broadcast Summing-up

Theorem. *The Spanning Tree Broadcast algorithm solves the broadcast problem in any connected synchronous network G when a rooted spanning tree T of G is known in advance. The time complexity of the algorithm (in rounds) is equal to the depth d of T . For the communication complexity, the algorithm transmits a total of $n - 1$ messages and the maximum size of a message is equal to the binary representation of the information to be broadcast.*

Broadcast without a given Spanning Tree

Problem:

- u_0 has some **information** it wishes to **send to all processors**
 - e.g., a **message** $\langle M \rangle$
 - additionally all nodes must have **terminated** in the end
- No spanning tree of the network G is given in advance
- The algorithm should also output a constructed spanning tree of G

Solution: Informal description

- All nodes **awake** initially
- If awake and have just received $\langle M \rangle$ from some neighbours,
 - Choose one of those neighbours as your **parent** and let him know
 - forward $\langle M \rangle$ to the rest of the neighbours
 - Wait for 1 round to collect children (if any) and then **sleep**
- If neighbours inform you that you are their parent,
 - add those processors to your **children** list
 - **sleep**
- If you are asleep, do nothing

Solution: Pseudocode

Algorithm Broadcast & Spanning tree construction

Code for processor u_i , $i \in \{0, 1, \dots, n - 1\}$:

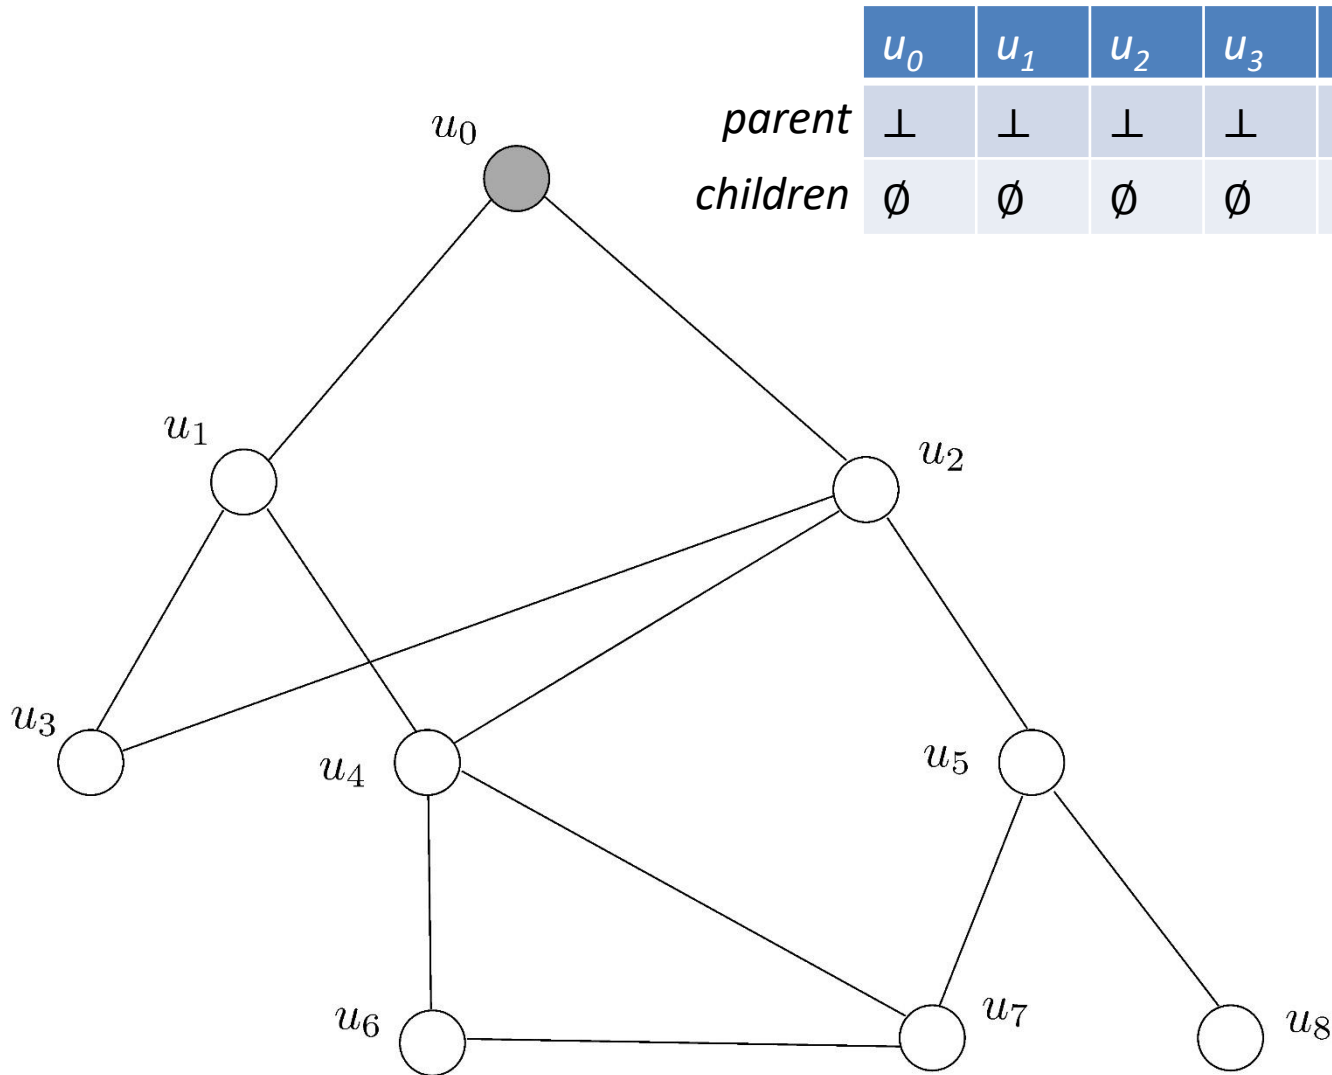
Initially $parent = \perp$ and $children = \emptyset$

if $u_i = u_0$ and $parent = \perp$ then // root has not yet sent $\langle M \rangle$
 send $\langle M \rangle$ to all neighbours
 $parent := u_i$

upon receiving $\langle M \rangle$ from neighbours N :
 if $parent = \perp$ then // u_i has not received $\langle M \rangle$ before
 $parent := u_j \in N$ // select one arbitrarily as parent
 send $\langle \text{"parent"} \rangle$ to u_j
 send $\langle M \rangle$ to all neighbours except those in N
 wait for one round to collect children if any and then terminate


upon receiving $\langle \text{"parent"} \rangle$ from neighbours N :
 add all $u_j \in N$ to $children$
 terminate

Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
<i>parent</i>	\perp	\perp	\perp	\perp	\perp	\perp	\perp	\perp	\perp
<i>children</i>	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

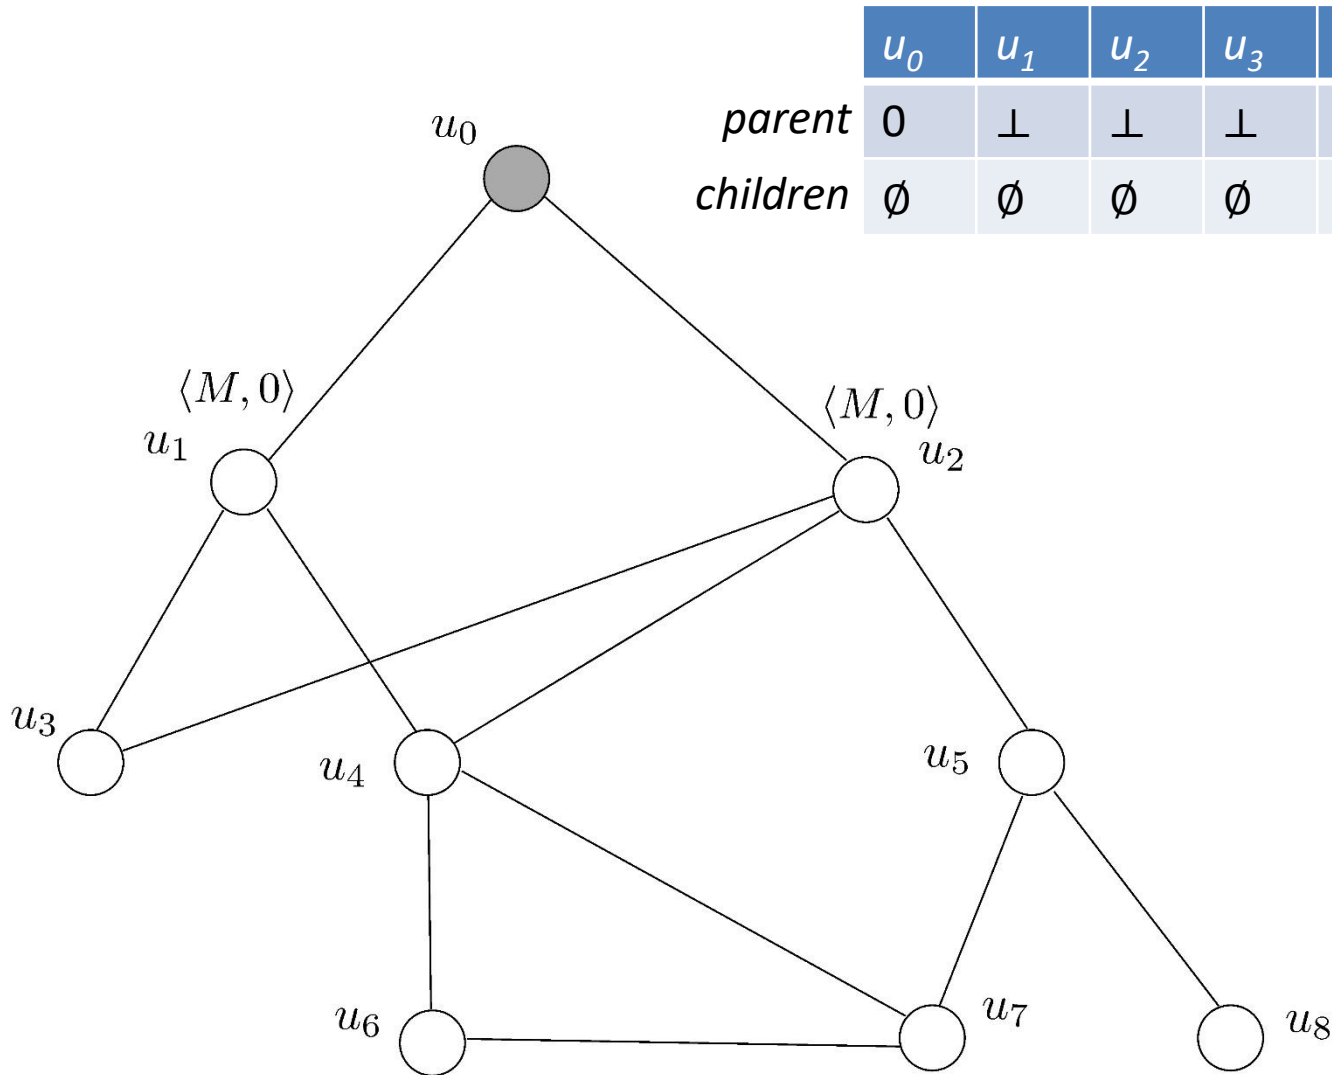
Example Execution



u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
0	\perp	\perp	\perp	\perp	\perp	\perp	\perp	\perp
\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

round = 1

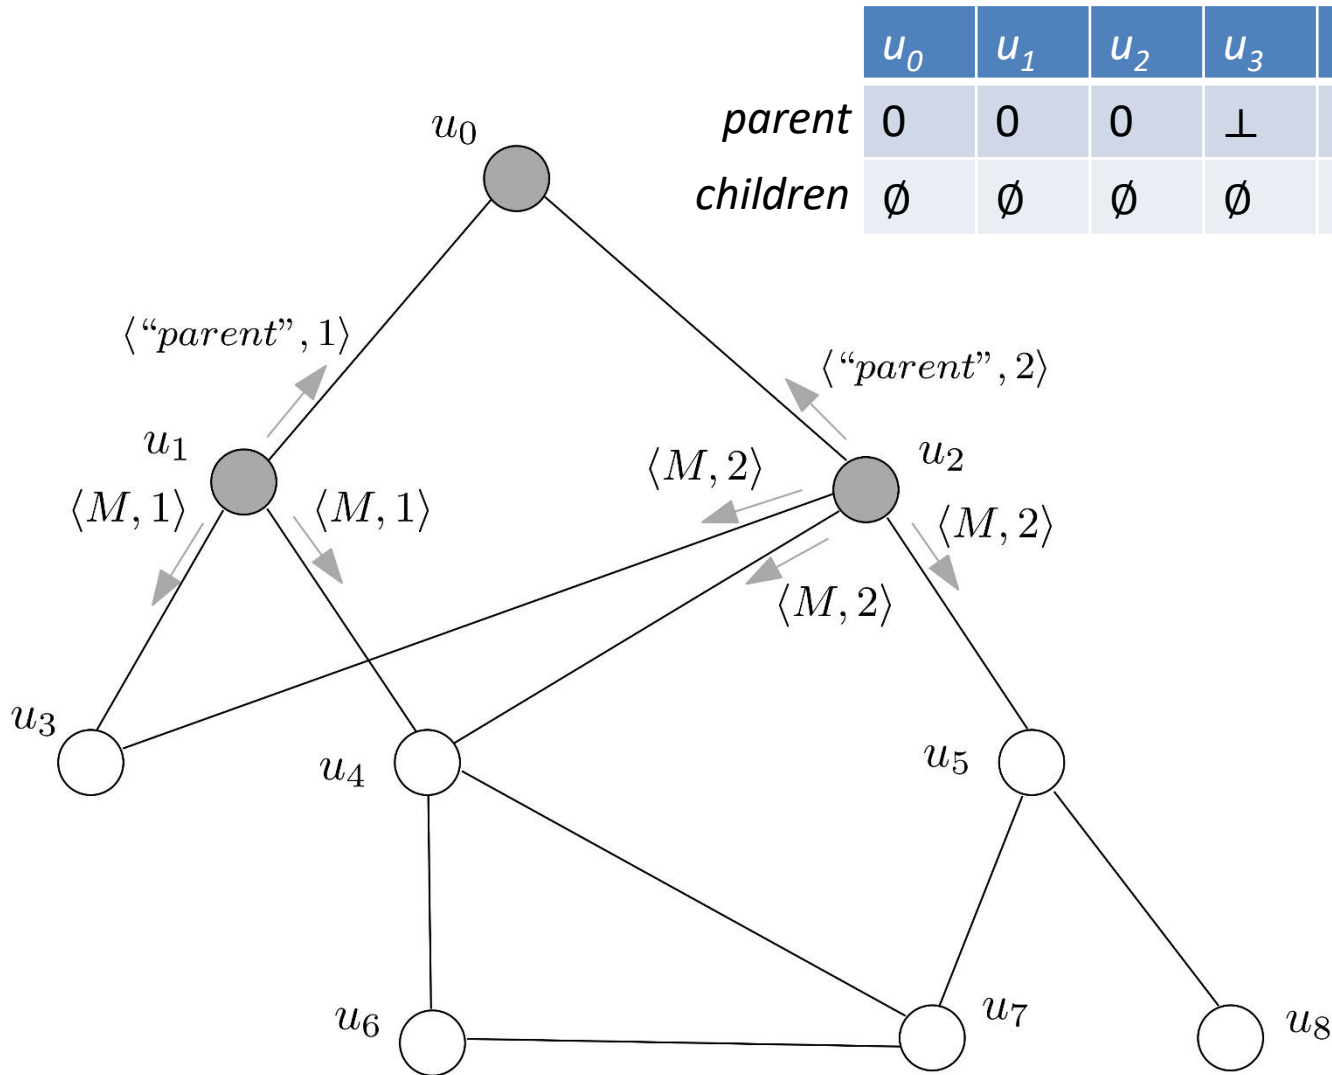
Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
<i>parent</i>	0	\perp	\perp	\perp	\perp	\perp	\perp	\perp	\perp
<i>children</i>	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

round = 1

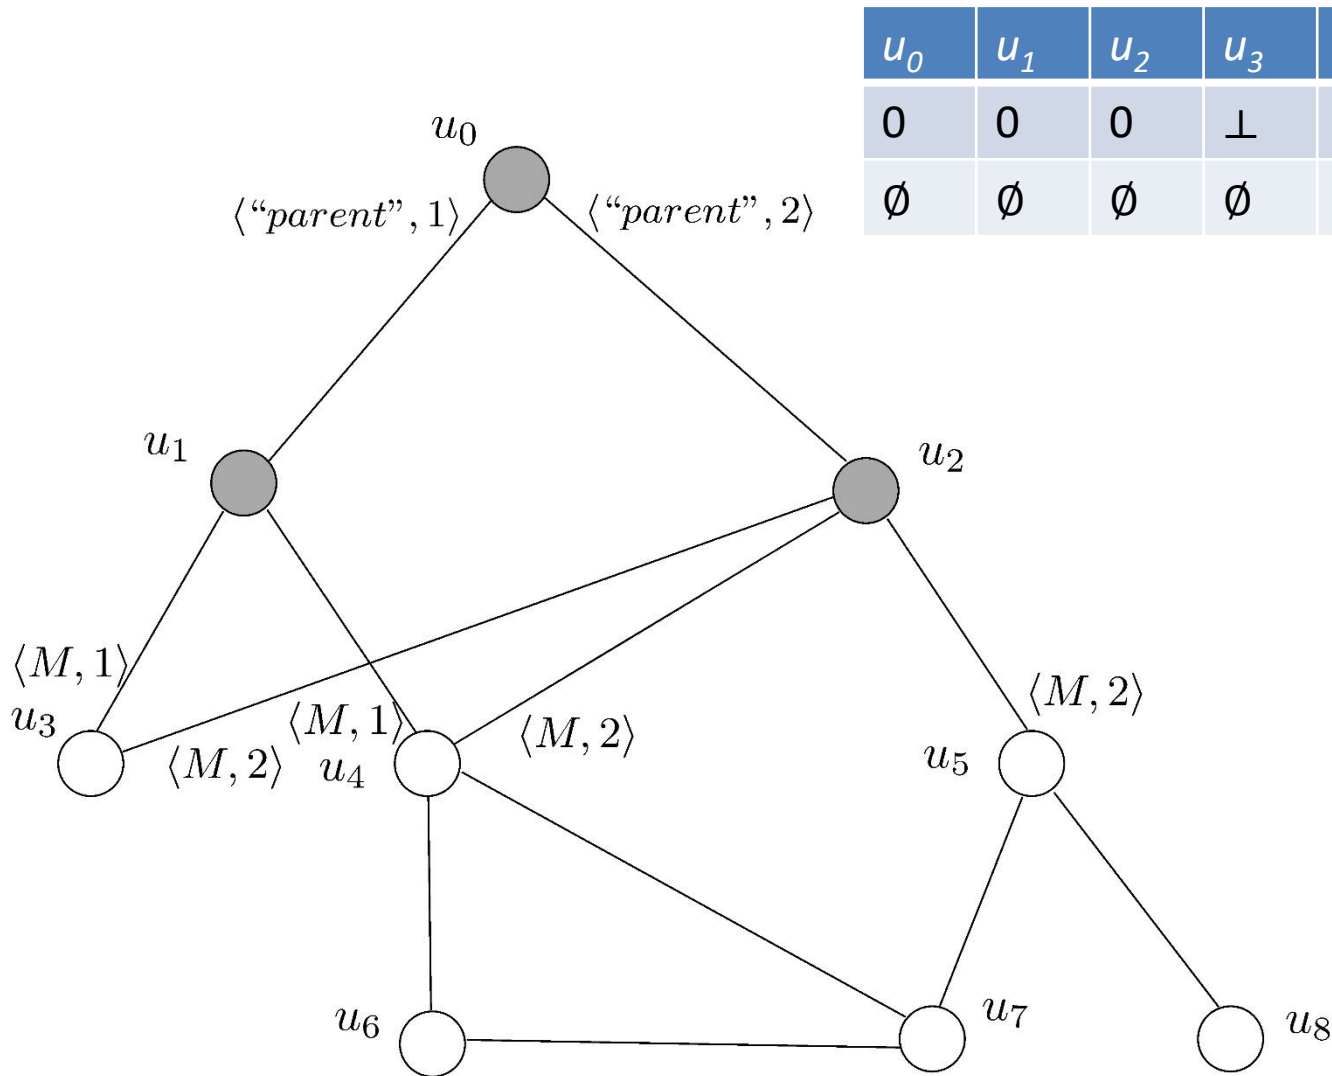
Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
<i>parent</i>	0	0	0	\perp	\perp	\perp	\perp	\perp	\perp
<i>children</i>	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

round = 2

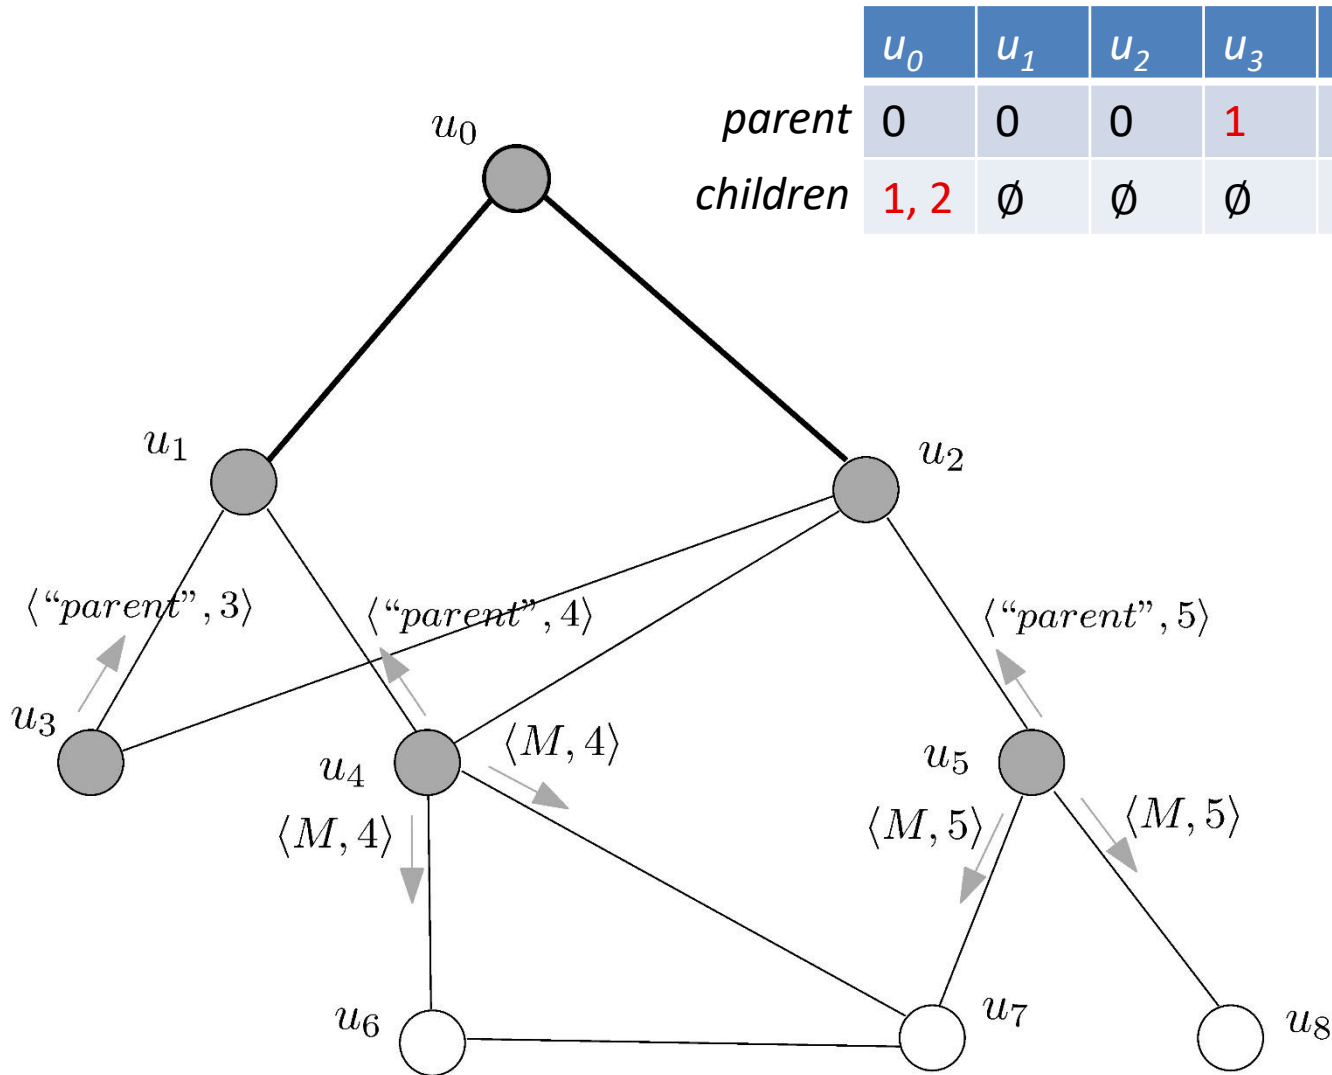
Example Execution



u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
0	0	0	\perp	\perp	\perp	\perp	\perp	\perp
\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

round = 2

Example Execution



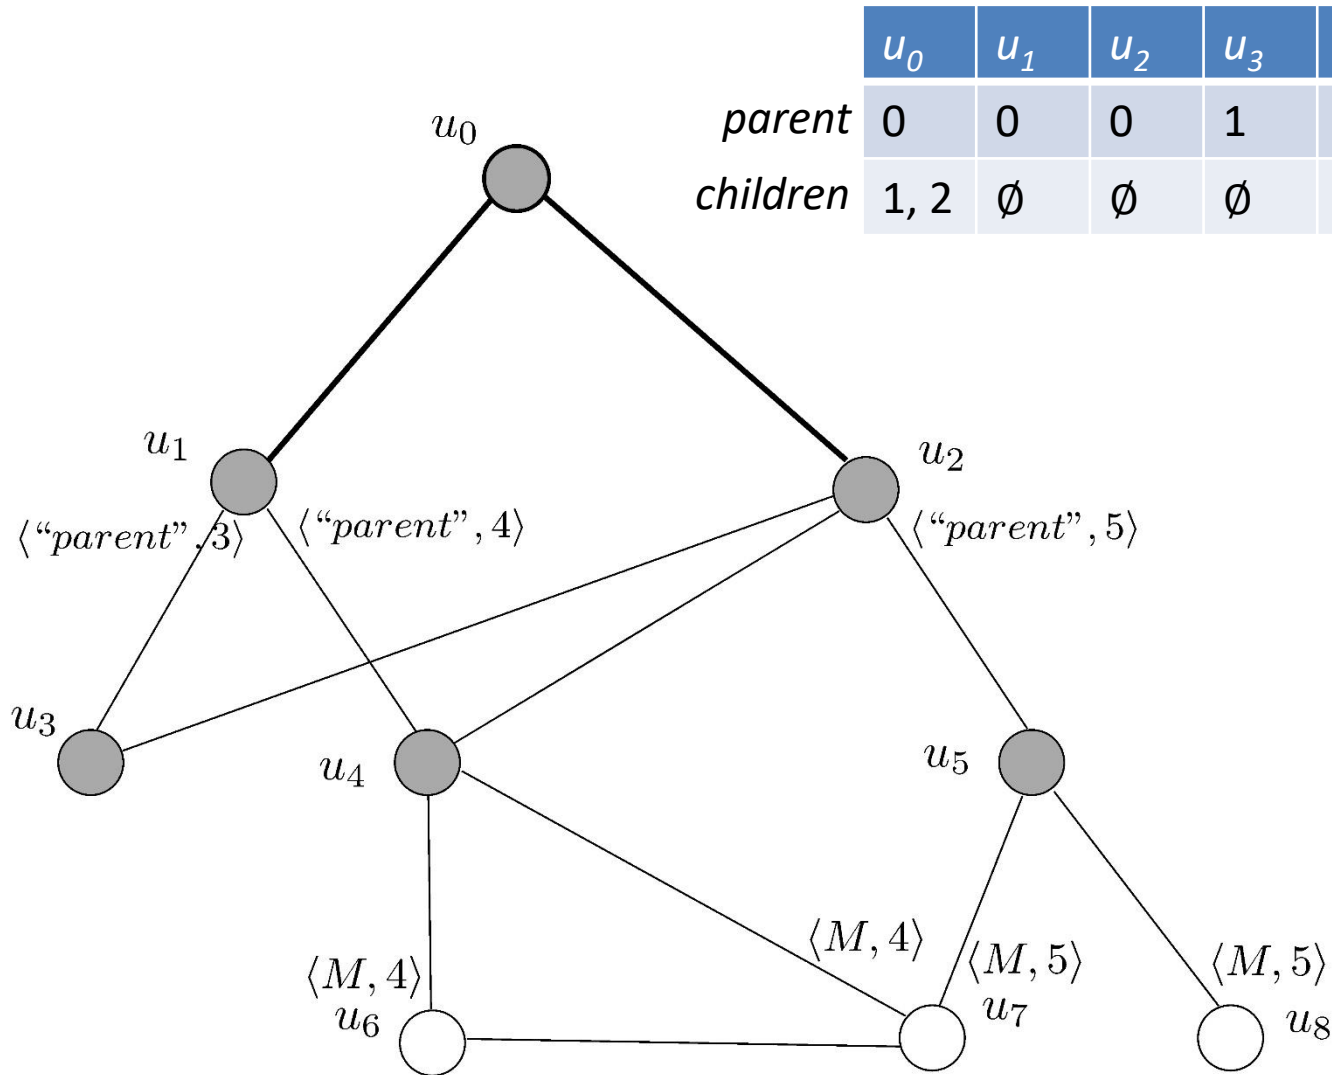
parent

children

u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
0	0	0	1	1	2	\perp	\perp	\perp
1, 2	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

round = 3


Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
<i>parent</i>	0	0	0	1	1	2	\perp	\perp	\perp
<i>children</i>	1, 2	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

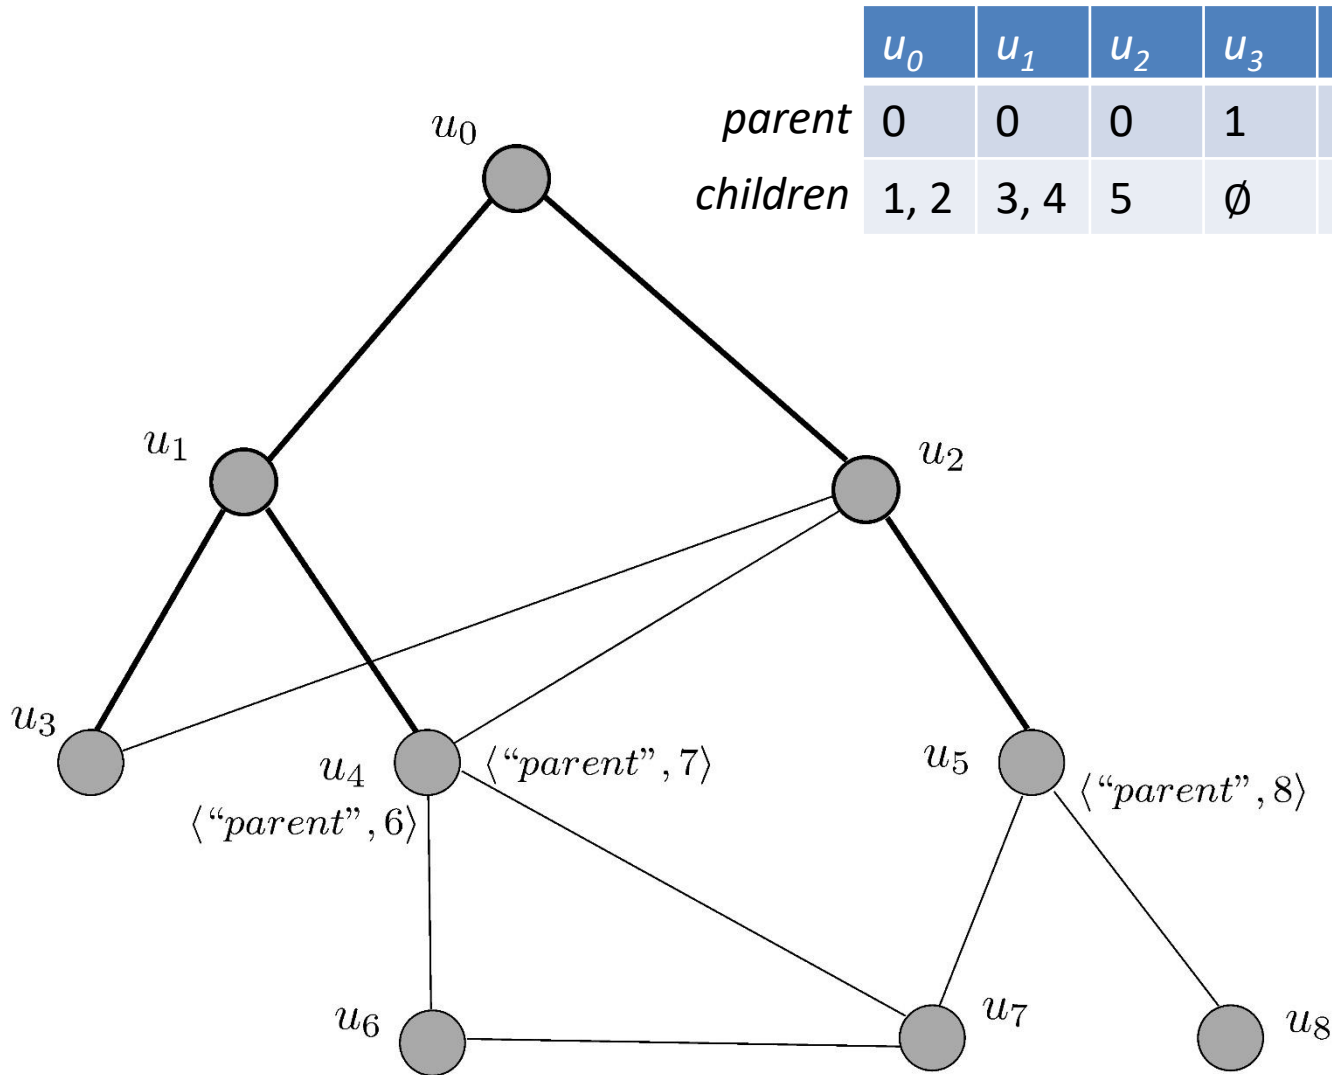
round = 3

Example Execution



round = 4

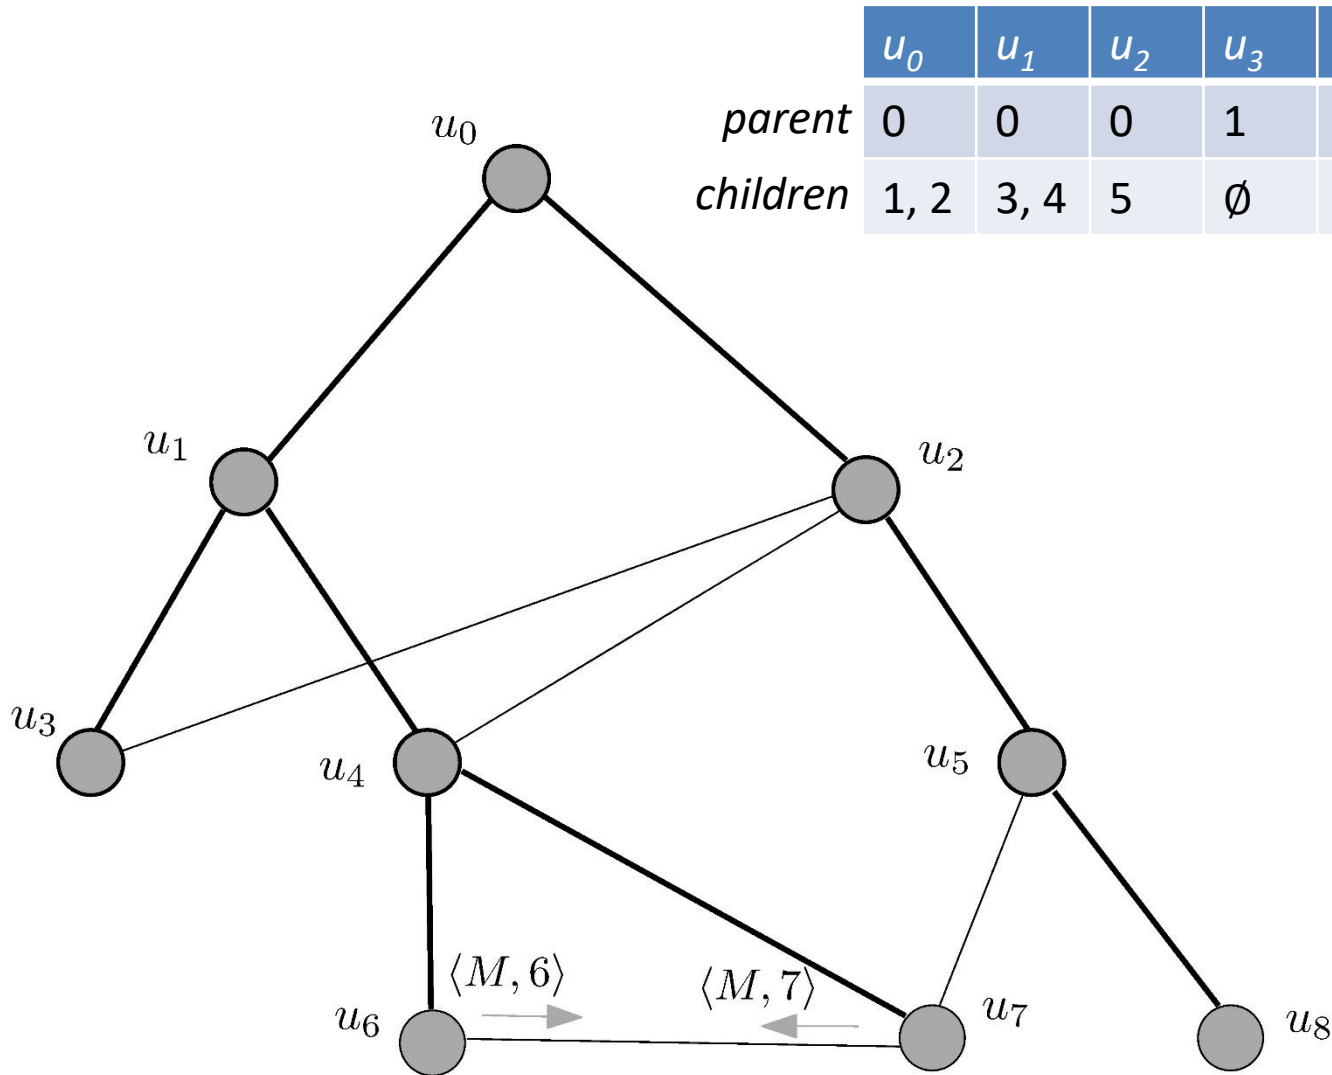
Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
<i>parent</i>	0	0	0	1	1	2	4	4	5
<i>children</i>	1, 2	3, 4	5	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset	\emptyset

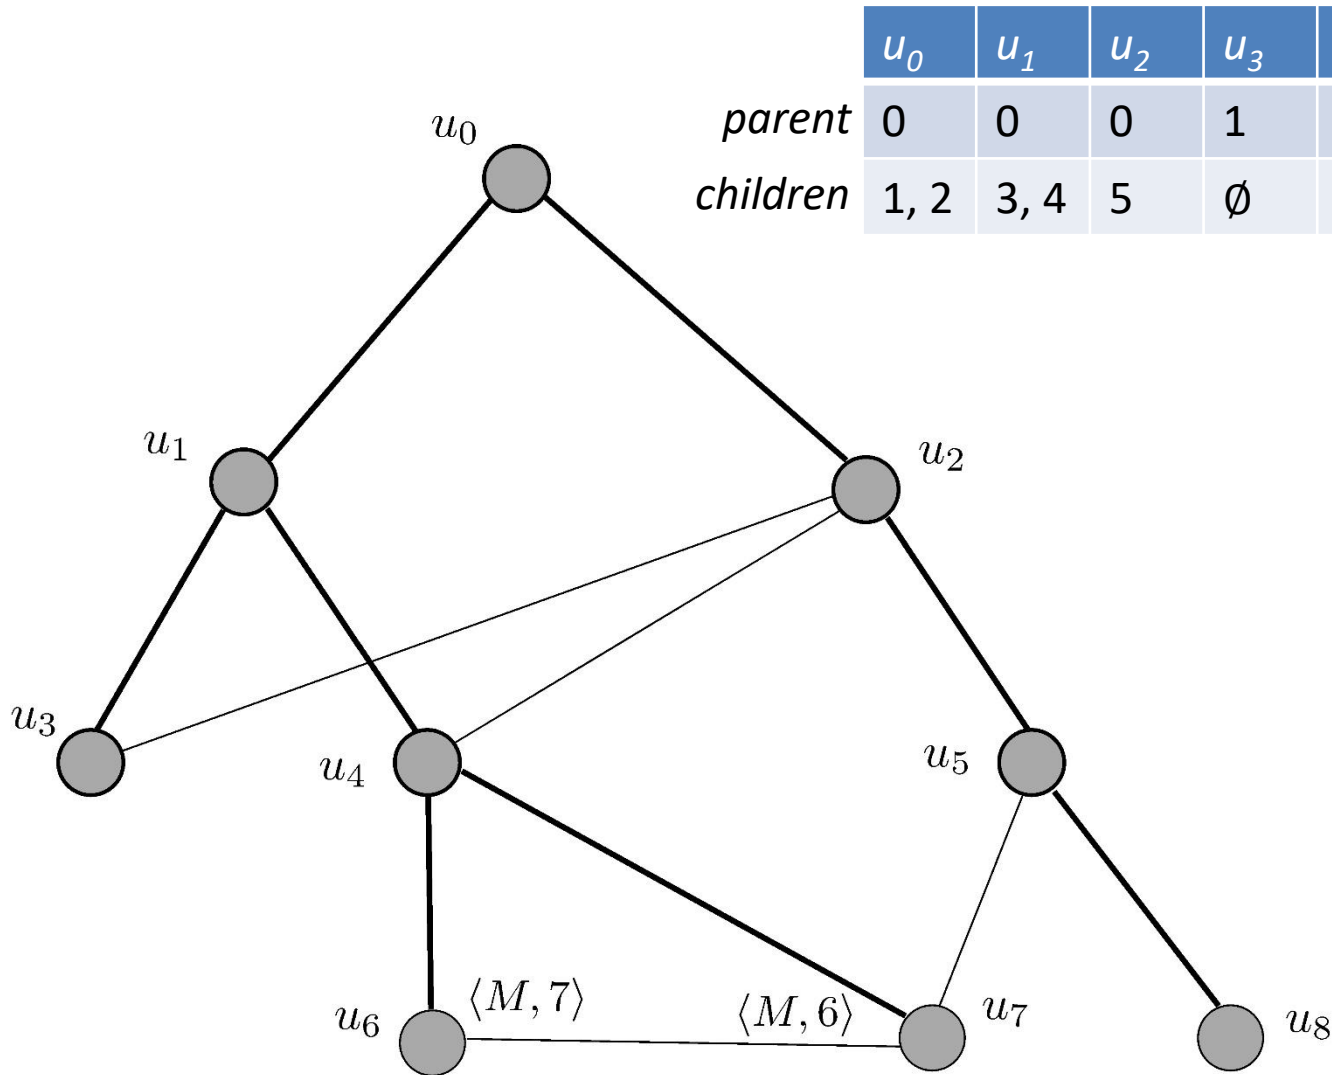
round = 4

Example Execution



round = 5

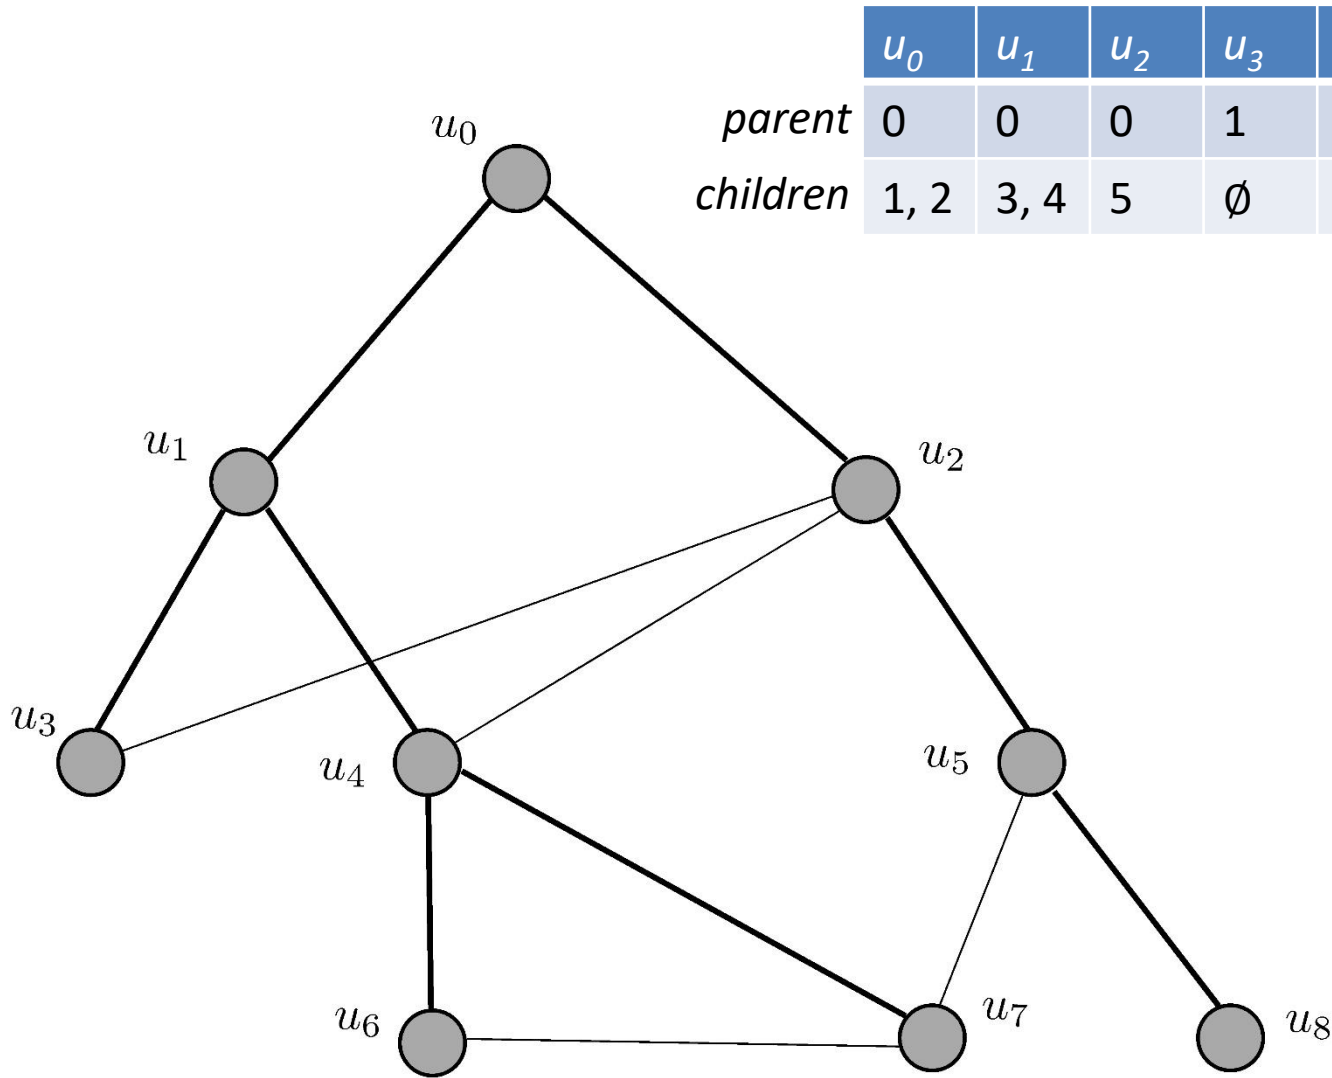
Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
parent	0	0	0	1	1	2	4	4	5
children	1, 2	3, 4	5	\emptyset	6, 7	8	\emptyset	\emptyset	\emptyset

round = 5


Example Execution



	u_0	u_1	u_2	u_3	u_4	u_5	u_6	u_7	u_8
<i>parent</i>	0	0	0	1	1	2	4	4	5
<i>children</i>	1, 2	3, 4	5	\emptyset	6, 7	8	\emptyset	\emptyset	\emptyset

round = 6

Example Execution



Correctness and Complexity

- **Correctness:**
 - correctness of broadcast
 - correctness of spanning tree construction
 - can also be shown that the constructed tree is always a **Breadth-first search (BFS) tree**
- **Time complexity:**
 - $O(D)$: where D is the **maximum distance** of a u_i from u_0 in G
- **Communication complexity:**
 - **size of messages**: sends **message $\langle M \rangle$** and an **id**
 - **$O(m)$ messages**: where m denotes the #edges of G

Can you prove these at home?