

Distributed Systems

COMP 212

Lecture 9

Othon Michail

Leader Election

Problem Statement

- *Elect a **unique leader processor** from among all the processors in the distributed system*
- Leader to be interpreted as:
 - **coordinator**
 - **master processor**
- Special case of **consensus/agreement**
- Processors should agree eventually on who they elect

Beyond Rings

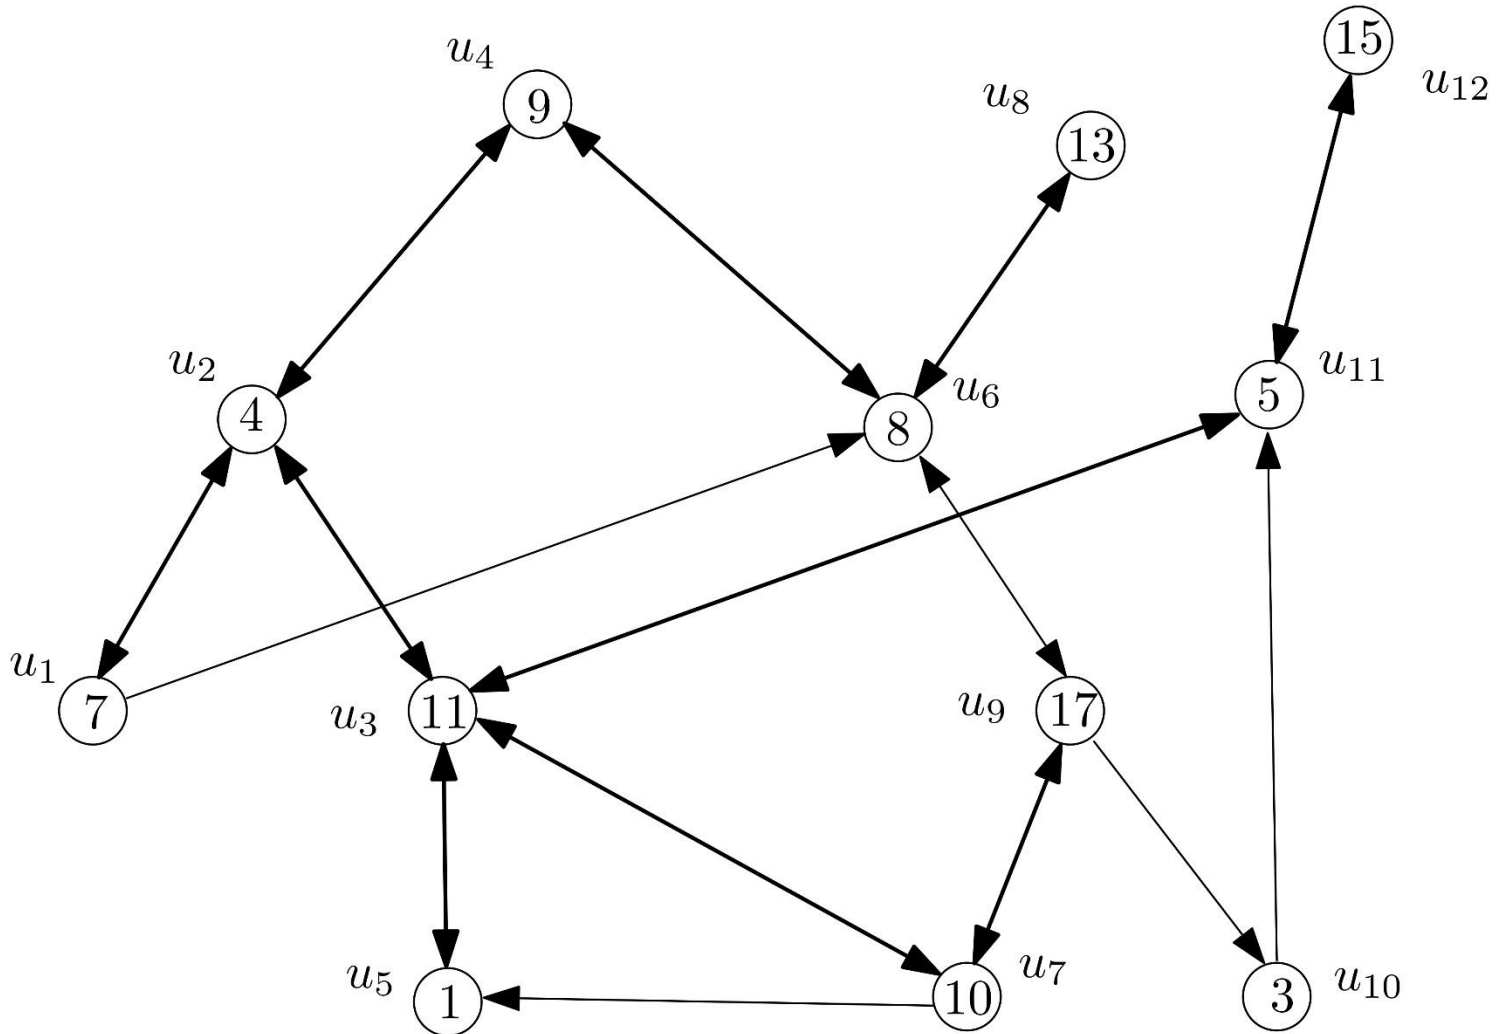
Leader Election in General Networks

Leader Election in General Networks

- Elect a *unique leader processor* from among all the processors in the distributed system
- Now the network can be **any strongly connected directed network**
 - **Strongly connected:** For every processors u, v there is a path from u to v and a path from v to u
 - e.g., the **directed ring** is just a special case
 - *Why can't we use LCR in this case?*
- Processors have **unique ids**

Leader Election in General Networks

- A strongly connected directed network



A Simple Algorithm based on Flooding

- Processors have **unique ids**, **do not know n** in advance, but **do know** the **diameter D** of the network
 - **Diameter:**
 - the **distance** between **two nodes** is given by the **shortest path between them**
 - Then the **diameter** of the network is determined by the **pair of nodes at maximum distance** (and is equal to that distance)
 - In other words, it is the **maximum shortest path in the network**
- **FloodMax algorithm:** solves the problem
- Uses **transmission**, **comparison**, and **storage** of ids
- **Main idea:** Flood the maximum id
 - LCR also does something like this but **does not require knowledge of D** and its termination condition **works only for rings**

FloodMax: Informal description

- All processors know the **diameter D** and their **own id** in advance
- All processors **remember** the **greatest id** that **they have “heard” so far** (initially their own)
- In every round all processors **send the greatest known to all their out-neighbours**
- After **D rounds** compare the largest heard to your own
 - if *greatest heard* = *own id*, **declare yourself the leader**
 - otherwise, **declare yourself non-leader**
- **Intuitively:**
 - The **maximum id** will manage to **reach the whole network**
 - **So everyone non-maximum will know** that there is a greater id and u_{max} can never receive a larger id

FloodMax: Pseudocode

Algorithm FloodMax

State of processor u_i :

- $myID_i$: holds the processor's unique id
- $maxID_i$: holds the greatest id “heard” so far
- $status_i \in \{\text{“unknown”, “leader”, “non-leader”}\}$:
indicates whether u_i has been elected (“leader”),
not elected (“non-leader”) or doesn't know yet
 (“unknown”)

FloodMax: Pseudocode

Algorithm FloodMax

Code for processor u_i , $i \in \{1, 2, \dots, n\}$:

Initially:

u_i knows its own unique id stored in $myID_i$

$maxID_i := myID_i$

$status_i := \text{"unknown"}$

Also has access to the current round and knows the diameter D

if $round = 1$ then

send $\langle maxID_i \rangle$ to all out-neighbours

else

upon receiving $\langle inIDs \rangle$ from in-neighbours

// one or more ids arriving from neighbours

$maxID_i := \max(\{maxID_i\} \cup inIDs)$

// remember only the maximum "heard" so far

if $round \leq D$ then // $1 < round \leq D$

send $\langle maxID_i \rangle$ to all out-neighbours

else // $round = D + 1$

if $maxID_i = myID_i$ then // if equal to your own, no greater id exists in the network

$status_i := \text{"leader"}$ // therefore, elect yourself a leader


else // greater than own

$status_i := \text{"non-leader"}$ // therefore, declare yourself a non-leader


// observe that in the end all processors know the id of the elected leader, stored in their $maxID_i$

// variable

Example Execution




Example Execution




round = 1

Example Execution




round = 1

Example Execution




round = 2

Example Execution




round = 2

Example Execution




round = 2

Example Execution




round = 3

Example Execution




round = 3

Example Execution




round = 3

Example Execution




round = 4

Example Execution




round = 4

Example Execution



round = 4

Example Execution



round = 5

Correctness and Complexity

- **Correctness:**
 - exactly one processor is elected in the last round
- **Time complexity:**
 - $D + 1$ rounds (or D depending on the round model)
- **Communication complexity:**
 - size of messages: encoding in bits of the maximum id
 - $D \cdot m$ messages always
 - m is the number of directed links in the network

FloodMax Correctness

- We have to show that:
 - Exactly one processor u_i sets $status_i := \text{"leader"}$ in the last round
 - We know that the algorithm aims to elect the processor with the maximum id
 - Call it u_{max}
- Suffices to show that:
 - u_{max} outputs *“leader”* in the last round
 - Every other u_i outputs *“non-leader”* in the last round

FloodMax Correctness

Lemma. u_{max} outputs “leader” in round $D + 1$.

Proof. Trivial.

- We know that id_{max} (i.e., the id of u_{max}) is the **greatest id** in the network
 - Therefore, **in every round** it will hold at u_{max} that
 - **$maxID = myID$**
 - as u_{max} will never hear an id greater than its own
 - So, this will also hold in round $D + 1$
 - Then **$maxID = myID$ evaluates to “true”** at u_{max}
 - Therefore, u_{max} outputs “leader” by setting $status := \text{“leader”}$
-
- Essentially, by induction on the number of rounds r , we show that **$maxID = myID$ holds for every r** at u_{max}

FloodMax Correctness

Lemma. Every processor u_i other than u_{max} outputs “non-leader” in round $D + 1$.

Proof. It suffices to show that by the beginning of round $D + 1$ every processor u_i has received id_{max} (i.e., the id of u_{max})

- Because then it must hold that in round $D + 1$
 - $maxID_i = id_{max} > myID_i$
 - and u_i will set $status_i := \text{“non-leader”}$
- We will prove that: In round r , any u_i at distance r from u_{max} receives id_{max}
 - By induction on r

FloodMax Correctness

Proof (continued).

- $r = 1$: Trivially, as u_{max} sends id_{max} to all its neighbours
- Assume it holds for any round $r - 1 \geq 1$
 - that is, assume that all nodes at distance from u_{max} receive id_{max} in round $r - 1$
- Then it must hold also for round r
 - Because all those nodes at distance $r - 1$, in round r set $maxID_i := id_{max}$
 - Therefore send id_{max} to all their neighbours
 - Implies that all nodes at distance r , receive id_{max} in round r

□

FloodMax Correctness

Theorem. The FloodMax algorithm solves the leader election problem in any **strongly connected directed network** (provided the availability of unique ids and knowledge of the diameter D).

- Observe that the diameter D concerns the maximum distance in the whole network
- *Can you think of a more precise parameter to replace D in this algorithm?*
 - In the worst case it will be equal to D but
 - In other cases it may be less

FloodMax Correctness

- Observe that the diameter D concerns the maximum distance in the whole network
- *Can you think of a more precise parameter to replace D in this algorithm?*
 - It is the maximum distance from u_{max} to any other processor
 - known as the eccentricity of that node
 - Observe that it would be a bit artificial to assume that the algorithm knows this in advance
 - Knowing the diameter of the network as a whole is quite natural to assume

FloodMax Time Complexity

- $D + 1$ rounds
- *Can you see why?*

FloodMax Time Complexity

- $D + 1$ rounds
 - All processors perform a check in round $D + 1$
 - From that point on they do nothing
 - We could have explicitly added a **halt** (or **terminate**) command at that point
 - At that point one processor has been elected and all other know that they have not been elected
- Question: *Do they also know who the elected one is?*

FloodMax Time Complexity

- $D + 1$ rounds
 - All processors perform a check in round $D + 1$
 - From that point on they do nothing
 - We could have explicitly added a **halt** (or **terminate**) command at that point
 - At that point 1 has been elected and all other know that they have not been elected
- Question: *Do they also know who the elected one is?*
 - Yes: In their *maxID_i* variable

FloodMax Communication Complexity

- $D \cdot m$ messages always
 - m (also denoted $|E|$) is the number of directed links in the network
- *Can you see why?*


FloodMax Communication Complexity

- $D \cdot m$ messages always
 - m (also denoted $|E|$) is the number of directed links in the network
- In round $D + 1$ nothing is transmitted
 - Only local checks and termination decisions
- For the first D rounds though:
 - Every processor u_i sends $\max ID_i$ to all its out-neighbours
 - Therefore, every link has one message in every round transmitted through it


FloodMax Communication Complexity

- *Is this the **most message-efficient solution?***
- *Can you observe any “waste” of messages in FloodMax?*

Example Execution




Example Execution



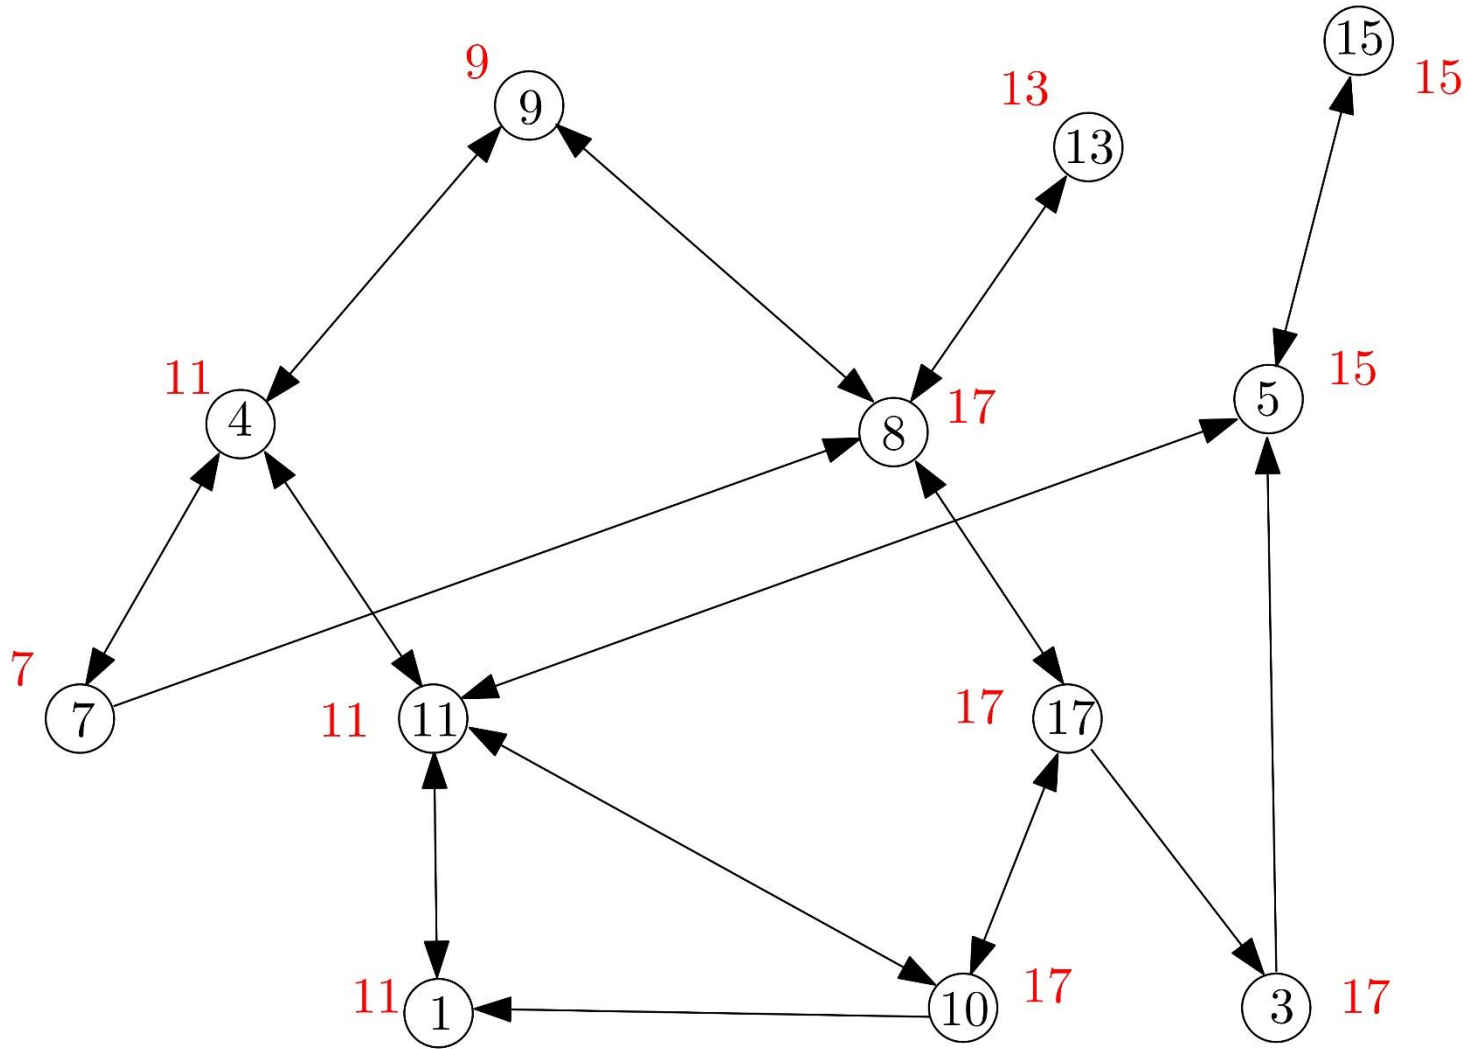
round = 1

Example Execution




round = 1

Example Execution




round = 2

Example Execution




Example Execution




round = 2

Example Execution



round = 3

Example Execution



OptFloodMax: An Improvement of FloodMax

- Same as FloodMax but now
 - Processors **do not** send their $maxID_i$ in every round
 - They **only** send it **whenever they hear a new maximum**
 - That is, only in **the** rounds that they **update their $maxID_i$**
- **Obvious** to see that it **reduces the #messages in various cases**
- **Not that obvious** yet whether it improves the **worst-case complexity**

OptFloodMax: Pseudocode

Algorithm OptFloodMax

Code for processor u_i , $i \in \{1, 2, \dots, n\}$:

Initially:

u_i knows its own unique id stored in $myID_i$

$maxID_i := myID_i$

$status_i := \text{"unknown"}$

$newInfo_i := \text{true}$ // an additional Boolean variable

Also has access to the current round and knows the diameter D



if $round = 1$ then

send $\langle maxID_i \rangle$ to all out-neighbours

else

upon receiving $\langle inIDs \rangle$ from in-neighbours

// one or more ids arriving from neighbours

if $\max(inIDs) > maxID_i$ then



$maxID_i := \max(\{maxID_i\} \cup inIDs)$

// remember only the maximum "heard" so far

$newInfo_i := \text{true}$



else



$newInfo_i := \text{false}$



if $round \leq D$ and $newInfo_i = \text{true}$ then // $1 < round \leq D$



send $\langle maxID_i \rangle$ to all out-neighbours

else if $round = D + 1$ then



if $maxID_i = myID_i$ then

// if equal to your own, no greater id exists in the network

$status_i := \text{"leader"}$

// therefore, elect yourself a leader

else

// greater than own

$status_i := \text{"non-leader"}$

// therefore, declare yourself a non-leader


// observe that in the end all processors know the id of the elected leader, stored in their $maxID_i$

// variable

OptFloodMax: Correctness and Complexity


- **Correctness:**
 - remains correct (needs proof)
- **Time complexity:**
 - same as in FloodMax (immediate)
- **Communication complexity:**
 - *We can show that it does not improve the worst-case complexity compared to FloodMax*
 - FloodMax: $D \cdot m = O(n^3)$ messages
 - We can show that in some cases also OptFloodMax transmits $\Theta(n^3)$ messages

OptFloodMax: Communication Complexity



- It happens that UIDs here are **1 through n (consecutive)** - **not necessary**
 - But their order in the network (combined with the specific structure of this network) is important for this result
- Remark: The network has **all** inverse links (to the left), not only the ones shown here

OptFloodMax: Communication Complexity



OptFloodMax: Communication Complexity

$$\begin{aligned}\#messages &= (n-1)^2 + \sum_{i=2} (n-i+1)^2 \\&= \left[\sum_{i=1}^n (n-i+1)^2 \right] - 2n + 1 \\&= \left[\sum_{i=1}^n i^2 \right] - 2n + 1 \\&= \frac{n^3}{3} + \frac{n^2}{2} + \frac{n}{6} - 2n + 1 \\&= \Theta(n^3)\end{aligned}$$

(Opt)FloodMax Further Improvement

- *Can you think of any additional improvement?*
- *Any waste of messages that still remains?*

(Opt)FloodMax Further Improvement

- *Can you think of any additional improvement?*
- *Any waste of messages that still remains?*
- **Yes: No need to send to a processor that just send you the new maximum**
 - Again won't improve the worst case complexity
 - Still, we gain something in many cases

Summary

- Leader election is crucial for distributed systems
 - breaks symmetry
 - allows for coordination
- If all processors are initially identical then
 - impossible to elect a leader even in very simple networks
 - e.g., a ring
- Adding unique ids breaks this inconvenient initial symmetry
- The LCR algorithm elects a leader in any ring network
 - simple conceptually, assumes unique ids
 - n rounds (or $2n$ for all to terminate), $O(n^2)$ messages
- The FloodMax algorithm elects a leader in any strongly connected network
 - like a generalisation of LCR
 - simple, assumes unique ids and knowledge of the diameter D
 - D rounds, $D \cdot m$ messages
- The OptFloodMax algorithm is an improvement of FloodMax
 - Decreases the number of messages in many cases
 - Does not improve the worst-case complexity
 - Still such improvements may be very important for real systems and applications where we are not always faced with the worst cases