

Distributed Systems

COMP 212

Lecture 15

Othon Michail

RPC/RMI vs Messaging

- RPC/RMI great in hiding communication in DSs
- But in some cases they are inappropriate
 - What happens if we cannot assume that the receiving side is “awake” and waiting to communicate?
 - Also, the default “synchronous, blocking” nature of RPC/RMI is often **too restrictive**

In such cases, something else is needed:

- **Messaging**
 - **Message-Oriented Middleware (MOM)**

Message-Oriented Communication and Streams

Some DS Comms. Terminology

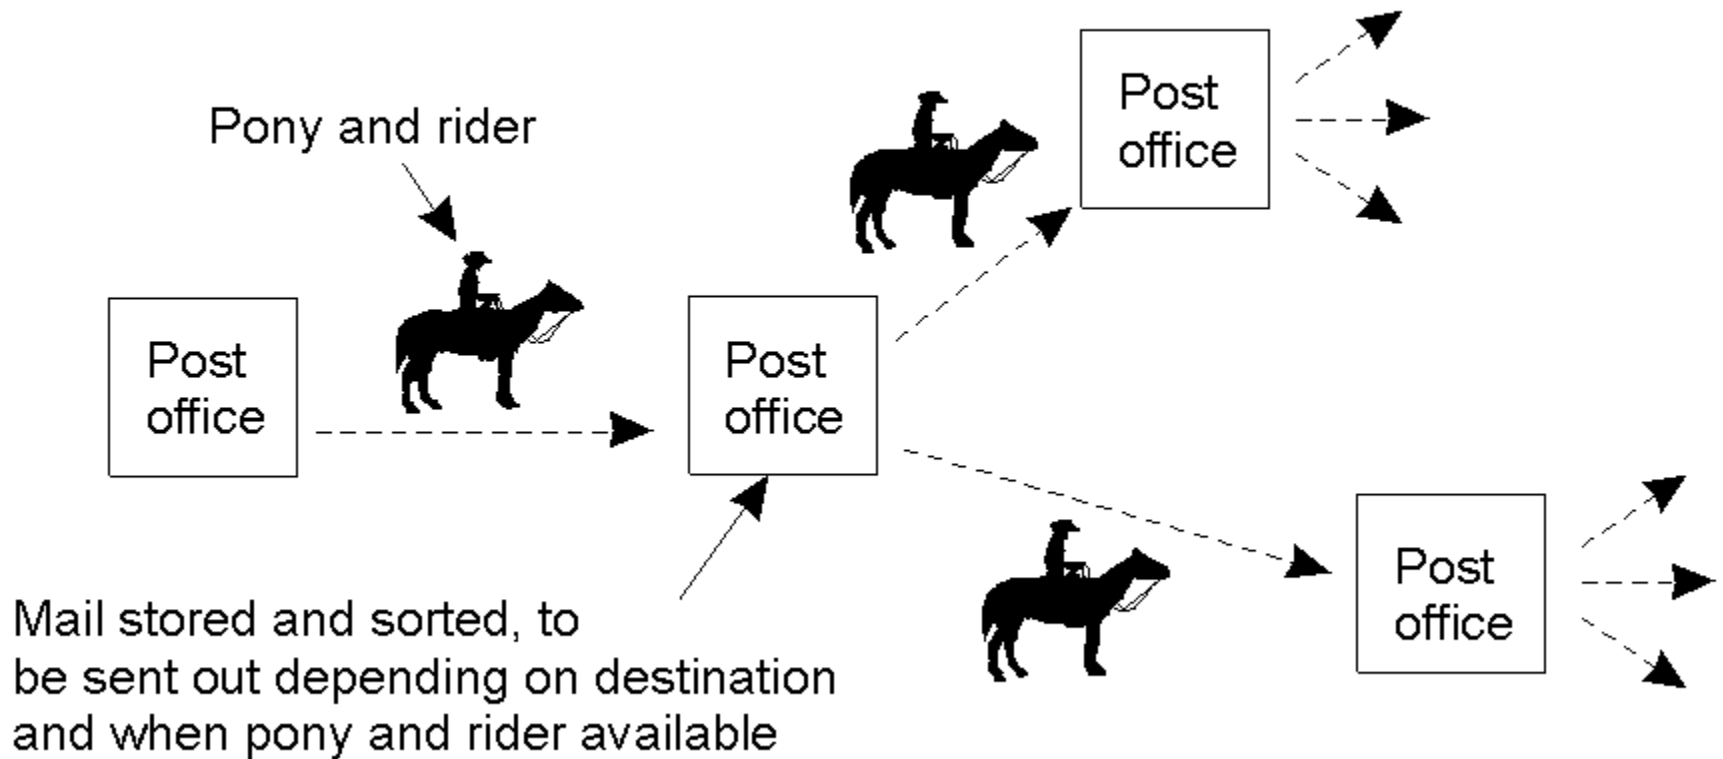
Persistent Communications:

- Once sent, the “sender” can stop executing. The “receiver” need not be operational at this time – the communications system **buffers** the message as required (until it can be delivered).
- [Can you think of an example?]

Contrast to Transient Communications:

- The message is only stored as long as the “sender” and “receiver” are executing. If problems occur, the message is simply **discarded** ...

Persistence and Synchronicity in Communication



- Persistent communication of letters back in the days of the Pony Express

More DS Comms. Terminology

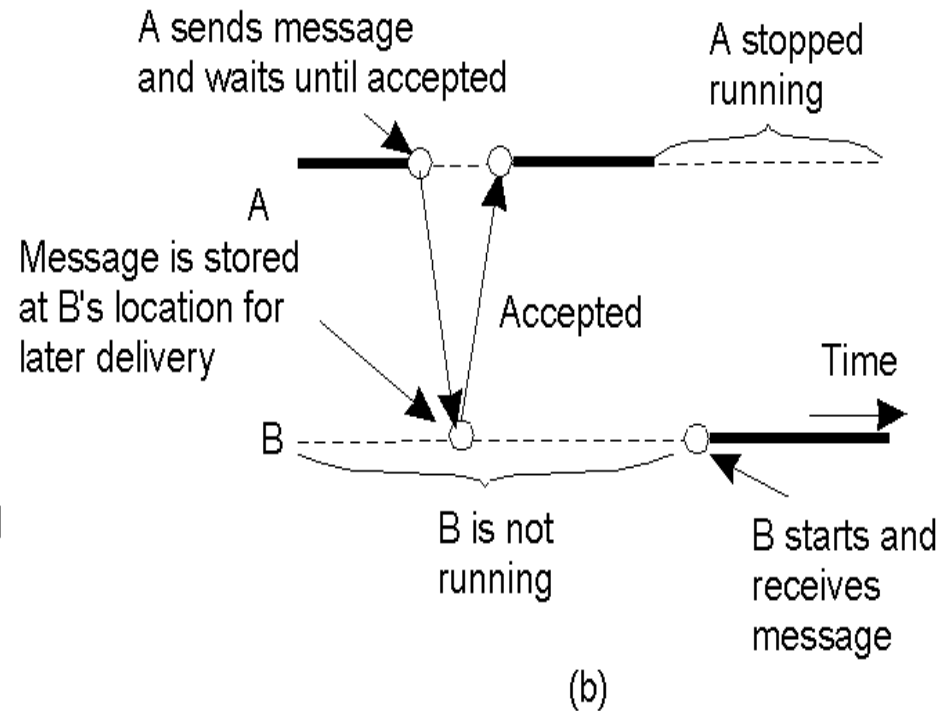
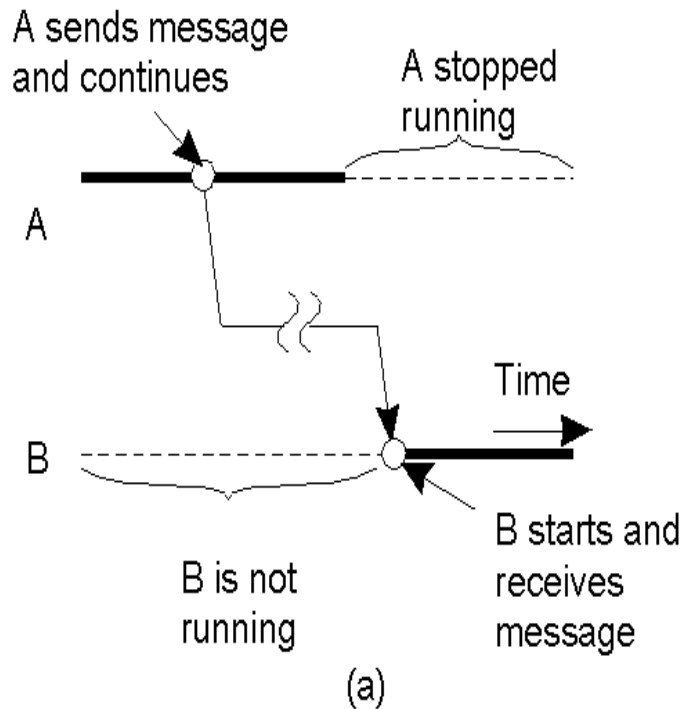
Asynchronous Communications:

- A sender **continues** with other work immediately upon sending a message to the receiver

Synchronous Communications:

- A sender **blocks, waiting** for a reply from the receiver before doing any other work
 - This tends to be the default model for RPC/RMI technologies
- We will have a look now at combinations

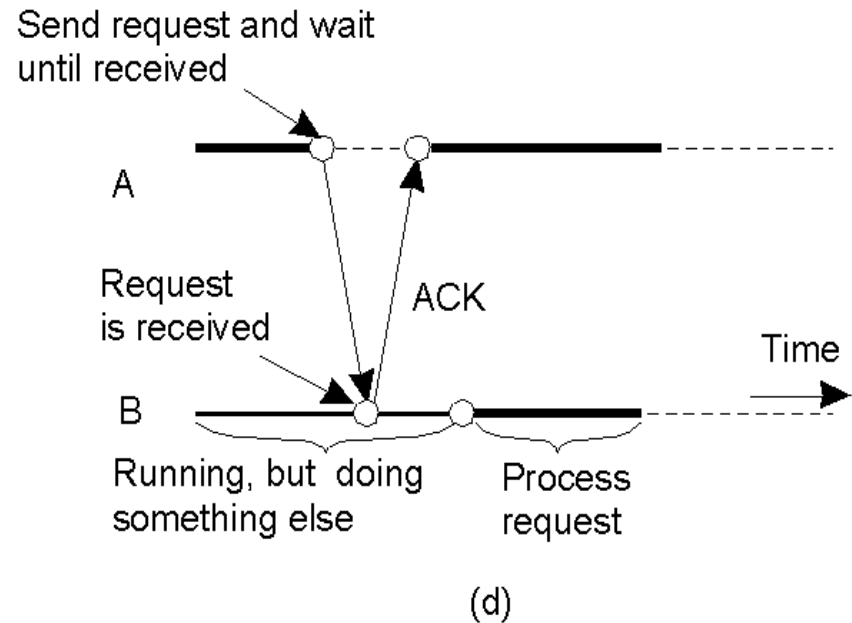
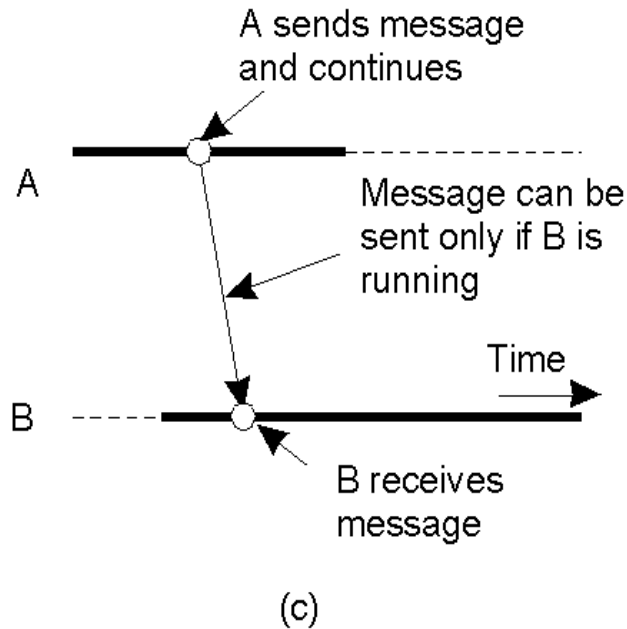
Classifying Distributed Communications (1)



a) **Persistent asynchronous** communication

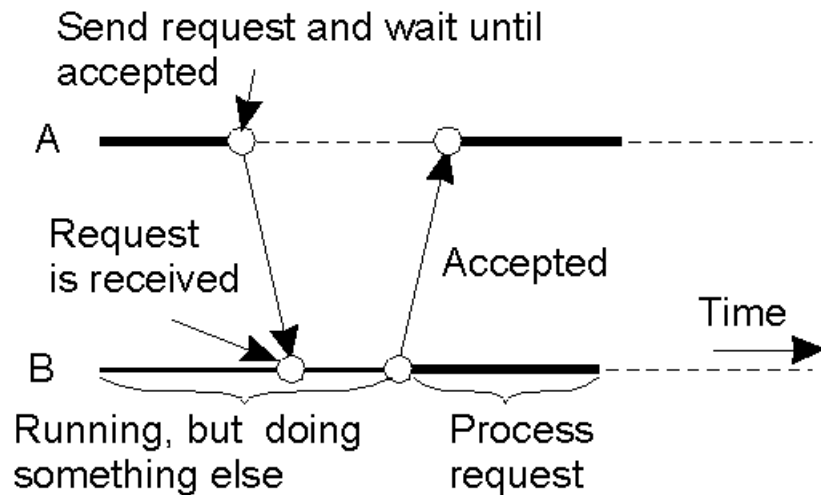
b) **Persistent synchronous** communication

Classifying Distributed Communications (2)

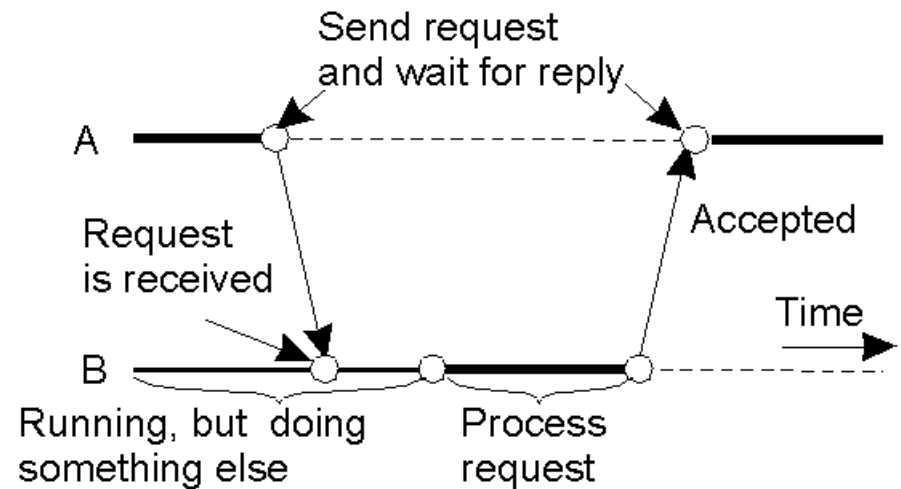


- c) Transient asynchronous communication
- d) Receipt-based transient synchronous communication

Classifying Distributed Communications (3)



(e)



(f)

- e) **Delivery-based transient synchronous** communication at message delivery
- f) **Response-based transient synchronous** communication

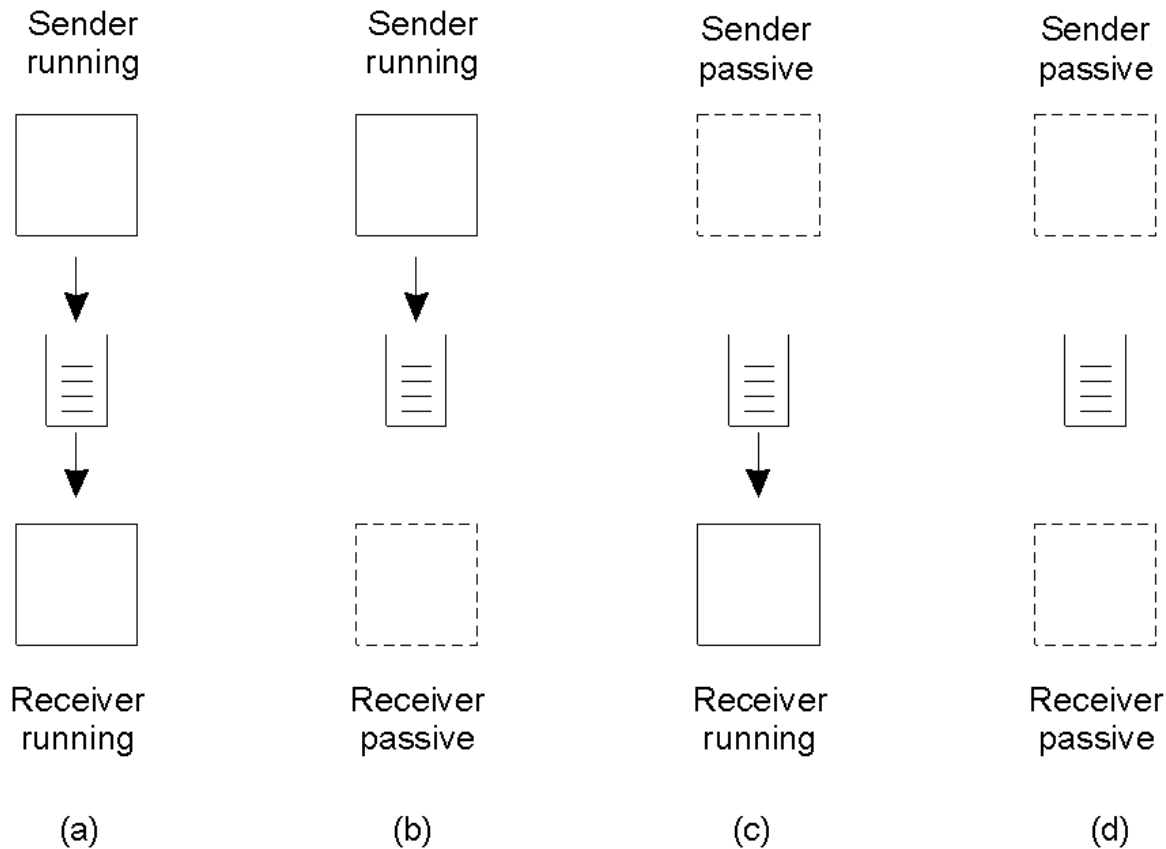
Example: Message-Passing Interface

- For high-performance **multicomputers**
- Specific network protocols (not TCP/IP)
- Message-based communication
- Primitives for all 4 forms of **transient** communication (+ variations)
- Over 100 functions
- Vendors
 - IBM, Intel, TMC, Meiko, Cray, Convex, Ncube

Message-Oriented Persistent Comms.

- Also known as: “message-queuing systems”
- They support persistent, asynchronous communications
- Typically, transport can take minutes (hours?) as opposed to seconds/milliseconds
- **The basic idea:** applications communicate by putting messages into and taking messages out of “message queues”
- **Only guarantee:** your message will eventually make it into the receiver’s message queue
- This leads to “loosely-coupled” communications

Message-Queuing Models



- Four combinations for “loosely-coupled” communications which use message-queues

Simple Interface

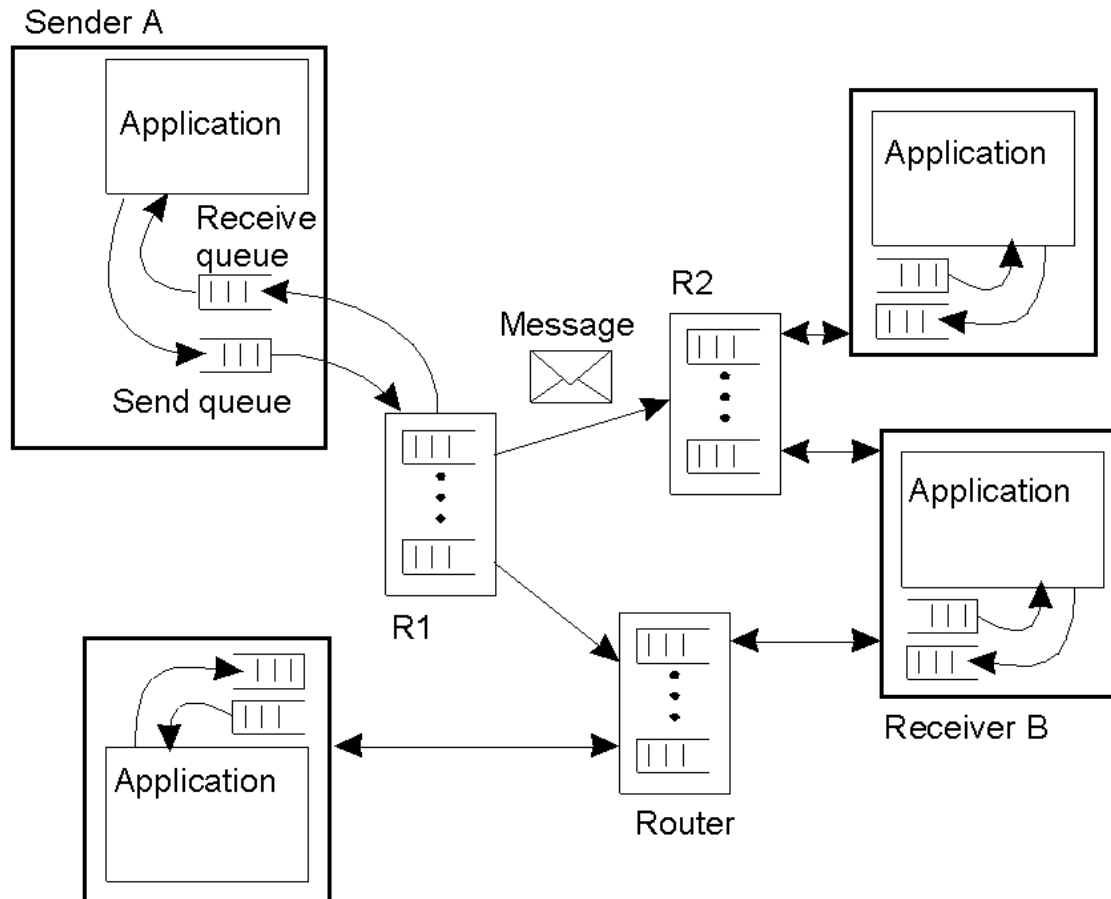
- Messages can contain any data
- Only important aspect: Must be properly addressed
 - Unique names of destination queues
- The system takes care of fragmenting and assembling large messages in a way transparent to applications
 - Extremely simple interface offered to applications

Primitive	Meaning
Put	Append a message to a specified queue
Get	Block until the specified queue is nonempty, and remove the first message
Poll	Check a specified queue for messages, and remove the first. Never block.
Notify	Install a handler to be called when a message is put into a specified queue

Message-Queuing System Architecture

- Messages are put into a **source queue**
- To be transferred to a **destination queue**
- Queues are distributed across multiple machines
 - The message-queuing system has to maintain a mapping of queues to network locations
 - Distributed database of **queue names** to **network locations**
 - e.g., email to somebody@liverpool.ac.uk
 - The mailing system queries DNS to find the network address of the recipient mail server to use for the actual message transfer
- A mechanism has to exist to move a message from a source queue to a destination queue
- This is the role of the **Queue Manager**
 - May interact **directly with** sending/receiving **applications**
 - May operate as **relays**, forwarding messages to other queue managers
 - May form a complete, application level, **overlay network**, on top of an existing computer network

General Architecture of a Message-Queuing System

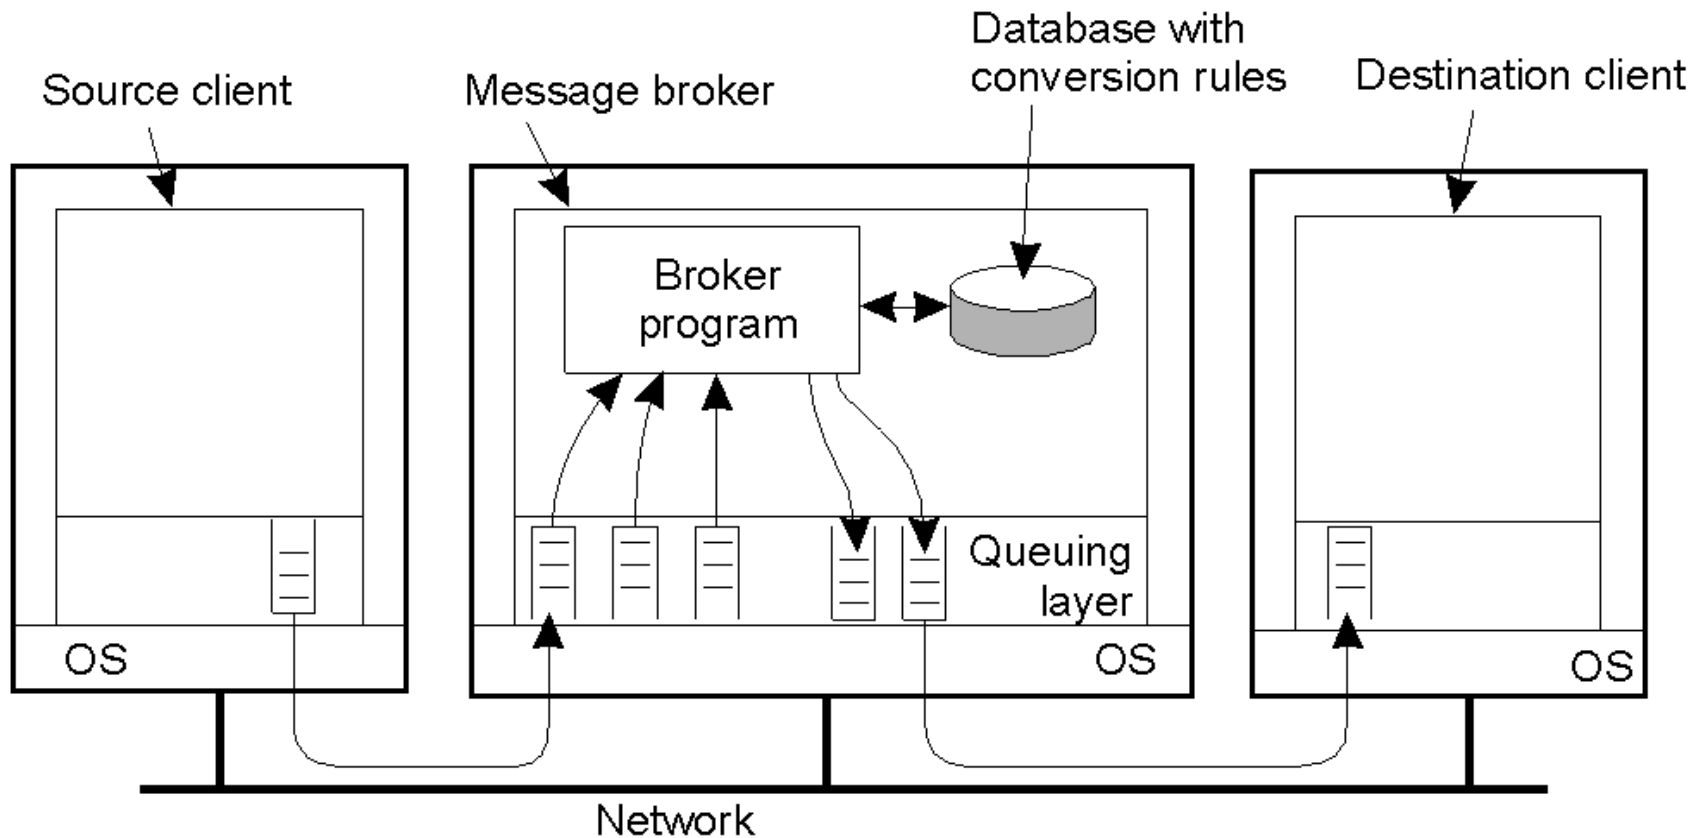


- The general organisation of a message-queuing system with routers
- The Queue Managers can reside within routers as well as within the DS end-systems

The Role of Message Brokers

- Often, there's a need to integrate new/existing apps into a “single, coherent Distributed Information System (DIS)”
 - In other words, it is not always possible to start with a **blank page**
 - distributed systems have to live in the real world
- **Problem:** different message formats exist in legacy systems (cooperation and adherence to open standards was not how things were done in the past)
- It may not be convenient to “force” legacy systems to adhere to a single, global message format (cost!?)
- It is often necessary to **learn to live with different formats**
- **How?**
 - By **Message Brokers**
 - Their role is to convert incoming messages in a format that can be understood by the destination application

Message Broker Organisation



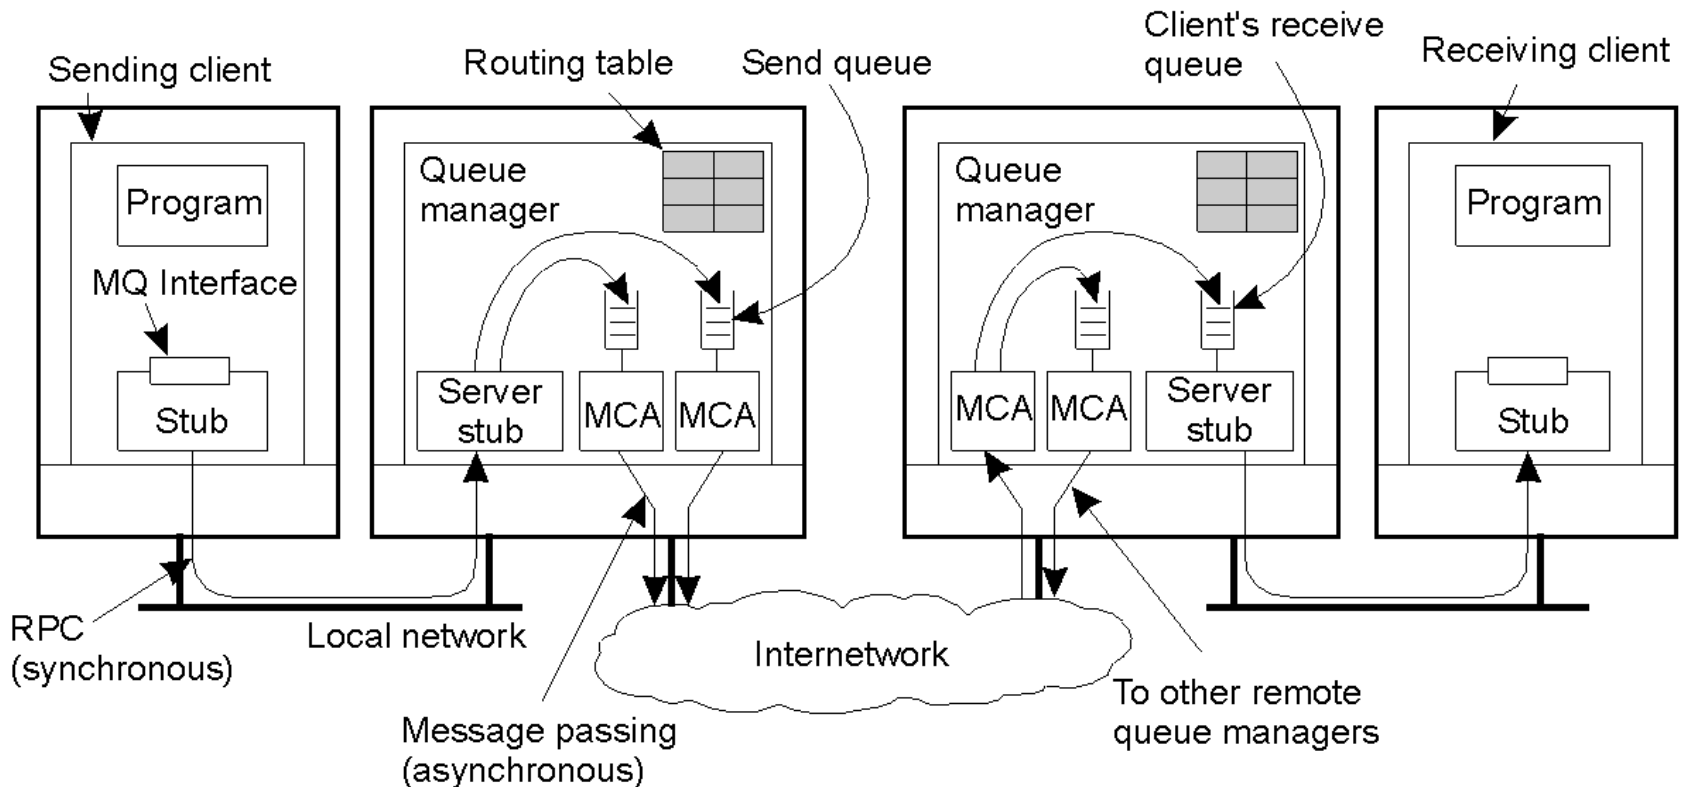
- The general organisation of a message broker in a message-queuing system
- Also known as an **interface engine**

Message-Queuing (MQ) Applications

- General-purpose MQ systems support a wide range of applications, including:
 - Electronic mail
 - Workflow
 - Groupware
 - Batch Processing
- **Most important MQ application area:**

The integration of a widely dispersed collection of **database applications** (which is all but impossible to do with traditional RPC/RMI techniques)

Example: IBM MQSeries



- General organisation of IBM's WebSphere MQSeries message-queuing system
 - Large-scale databases, finance

Stream-Oriented Communications

- With RPC, RMI and MOM, the effect that time has on correctness is of little consequence
- However, audio and video are **time-dependent data streams**
 - if the timing is off, the resulting “output” from the system will be incorrect
- Time-dependent information
 - known as **continuous media** communications

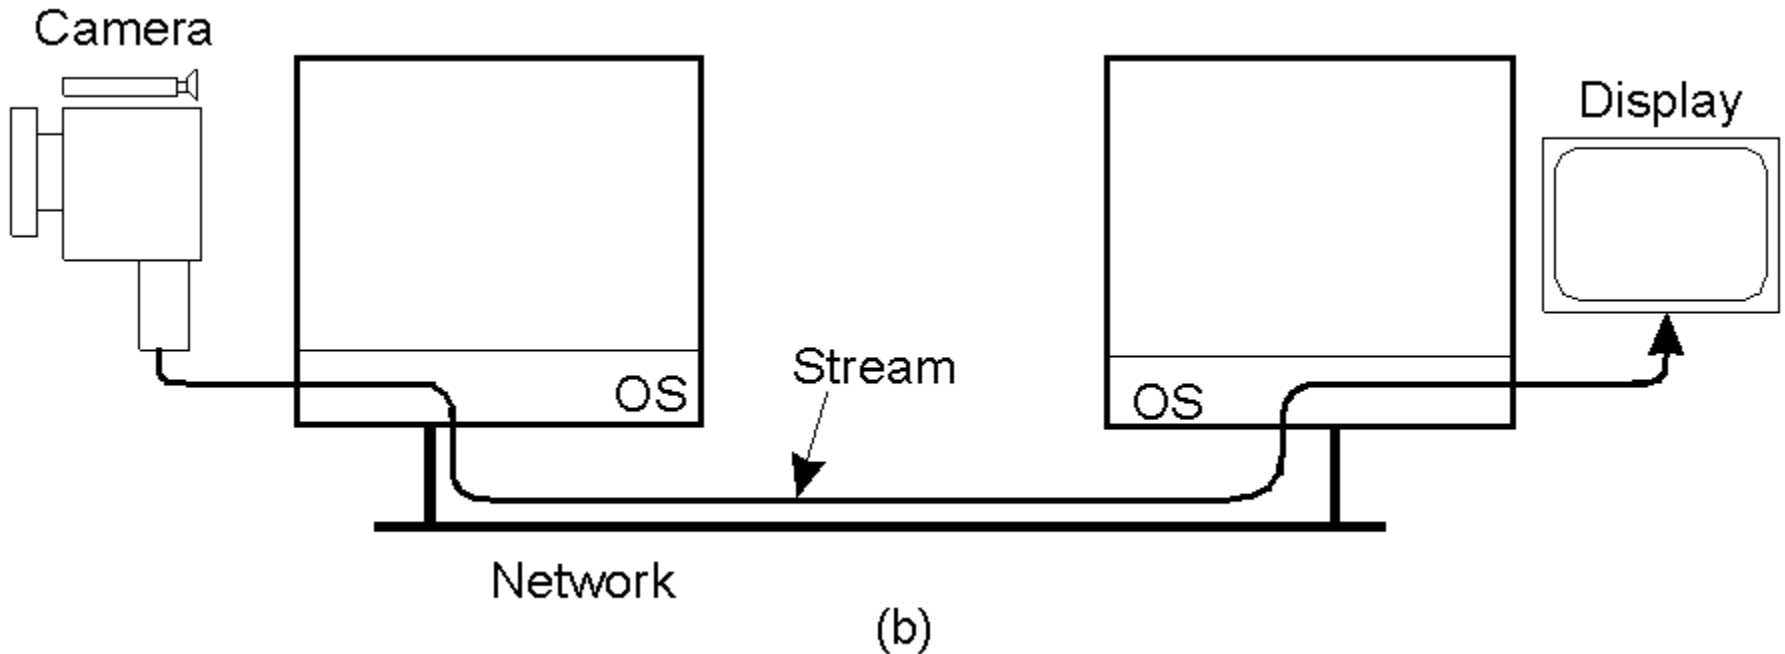
Examples:

- Audio: (PCM) samples played at 1/44100 sec intervals (exactly)
- Video: 30 frames per second (30-40 msec per image)
- **KEY MESSAGE:** Timing is crucial!

Transmission Modes

- **Asynchronous transmission mode**
 - The data stream is transmitted in order, but there's no timing constraints placed on the actual delivery (e.g., File Transfer)
- **Synchronous transmission mode**
 - The maximum end-to-end delay is defined (but data can travel faster)
- **Isochronous transmission mode** –
 - Data transferred **on time**
 - There's a **maximum** and **minimum** end-to-end delay (known as **bounded jitter**)
 - Known as **streams**
 - Very useful for multimedia systems

End-device to End-device Streams



- Setting up a stream **directly** between two devices
 - i.e., no inter-networked processes

Two Types of Streams


- **Simple Streams** – one single sequence of data, for example: voice
- **Complex Streams** – several sequences of data (sub-streams) that are time-related
 - Think of a movie:
 - 1 stream for video
 - 2 streams for stereo sound
 - 1 stream for subtitles, ...
- This leads to **data synchronisation problems** ... which are not at all easy to deal with

Quality of Service

- Definition: “Ensuring that the temporal relationships in the stream can be preserved”
- QoS is all about three things:
 1. Timeliness
 2. Volume
 3. Reliability
- Most current operating systems and networks **do not** include the QoS management facilities
- Bleeding edge of the discipline


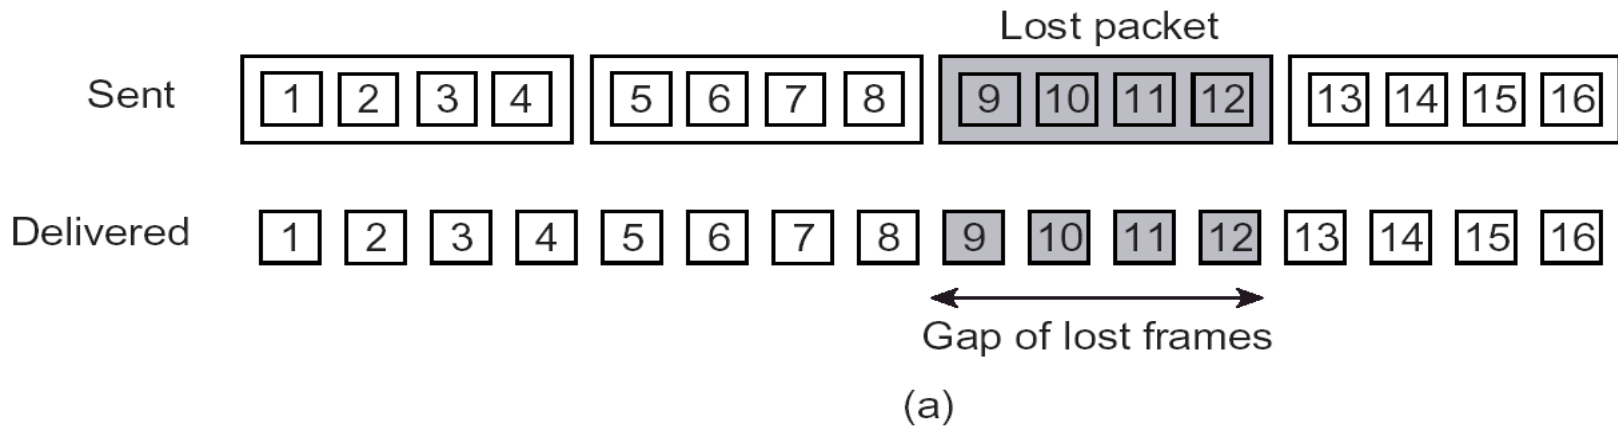
Enforcing QoS (1)

- Use **buffers** to reduce jitter
- **Note:** the larger the buffer the greater the **initial delay** of starting playback



Enforcing QoS (2)

- Underlying **best-effort service**: packets may be lost
- Reducing the effect of **packet loss**



Distributed Comms. - Summary

- Power and flexibility essential, as network programming primitives are **too primitive**
- **Middleware Comms. Mechanisms** – providing support for a **higher-level of abstraction**
- **RPC** and **RMI**: synchronised, transient
- **MOM**: convenient, asynchronous, persistent
- **Streams**: a special case, useful when dealing with **temporally-related data** (not easy)